

# Robot Autonomy

---

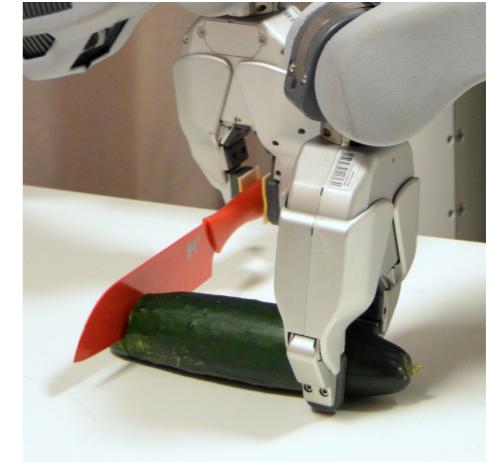
## Lecture 3: **Self-supervised Learning**

---

Oliver Kroemer

# Motivation

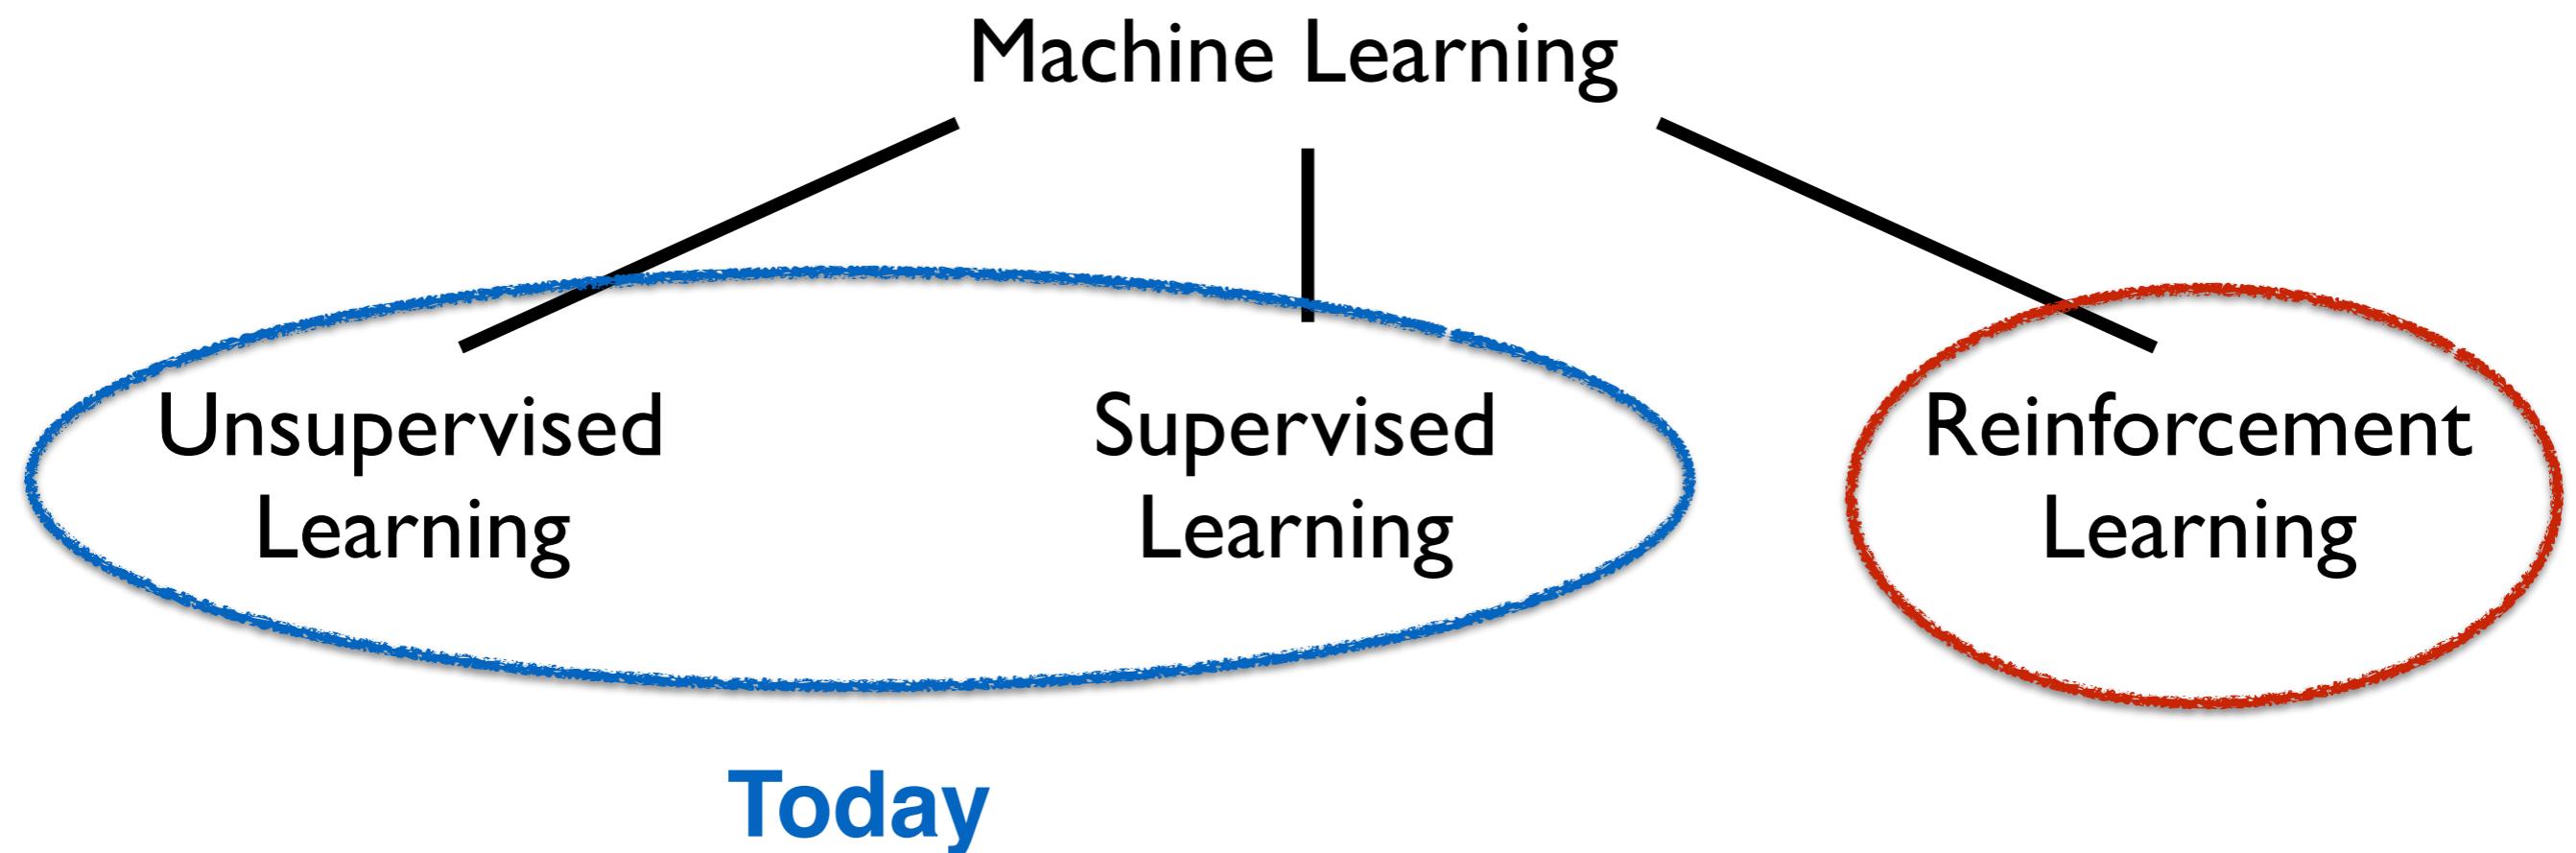
- Robots need to understand their environment



- ▶ Identify object properties and action possibilities
- ▶ Predict the effects of different interactions
- ▶ Use this information to perform tasks
- Robot to increase autonomy by learning on its own

# Machine Learning

---



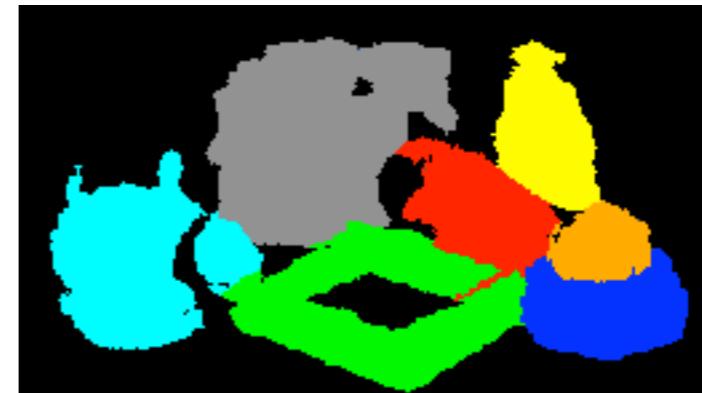
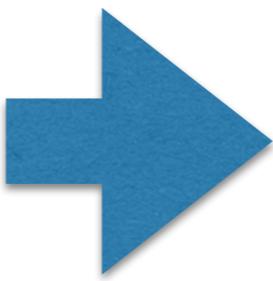
- Robots can acquire information from many sources
- Will be focusing on learning from own experiences

# Interactive Perception

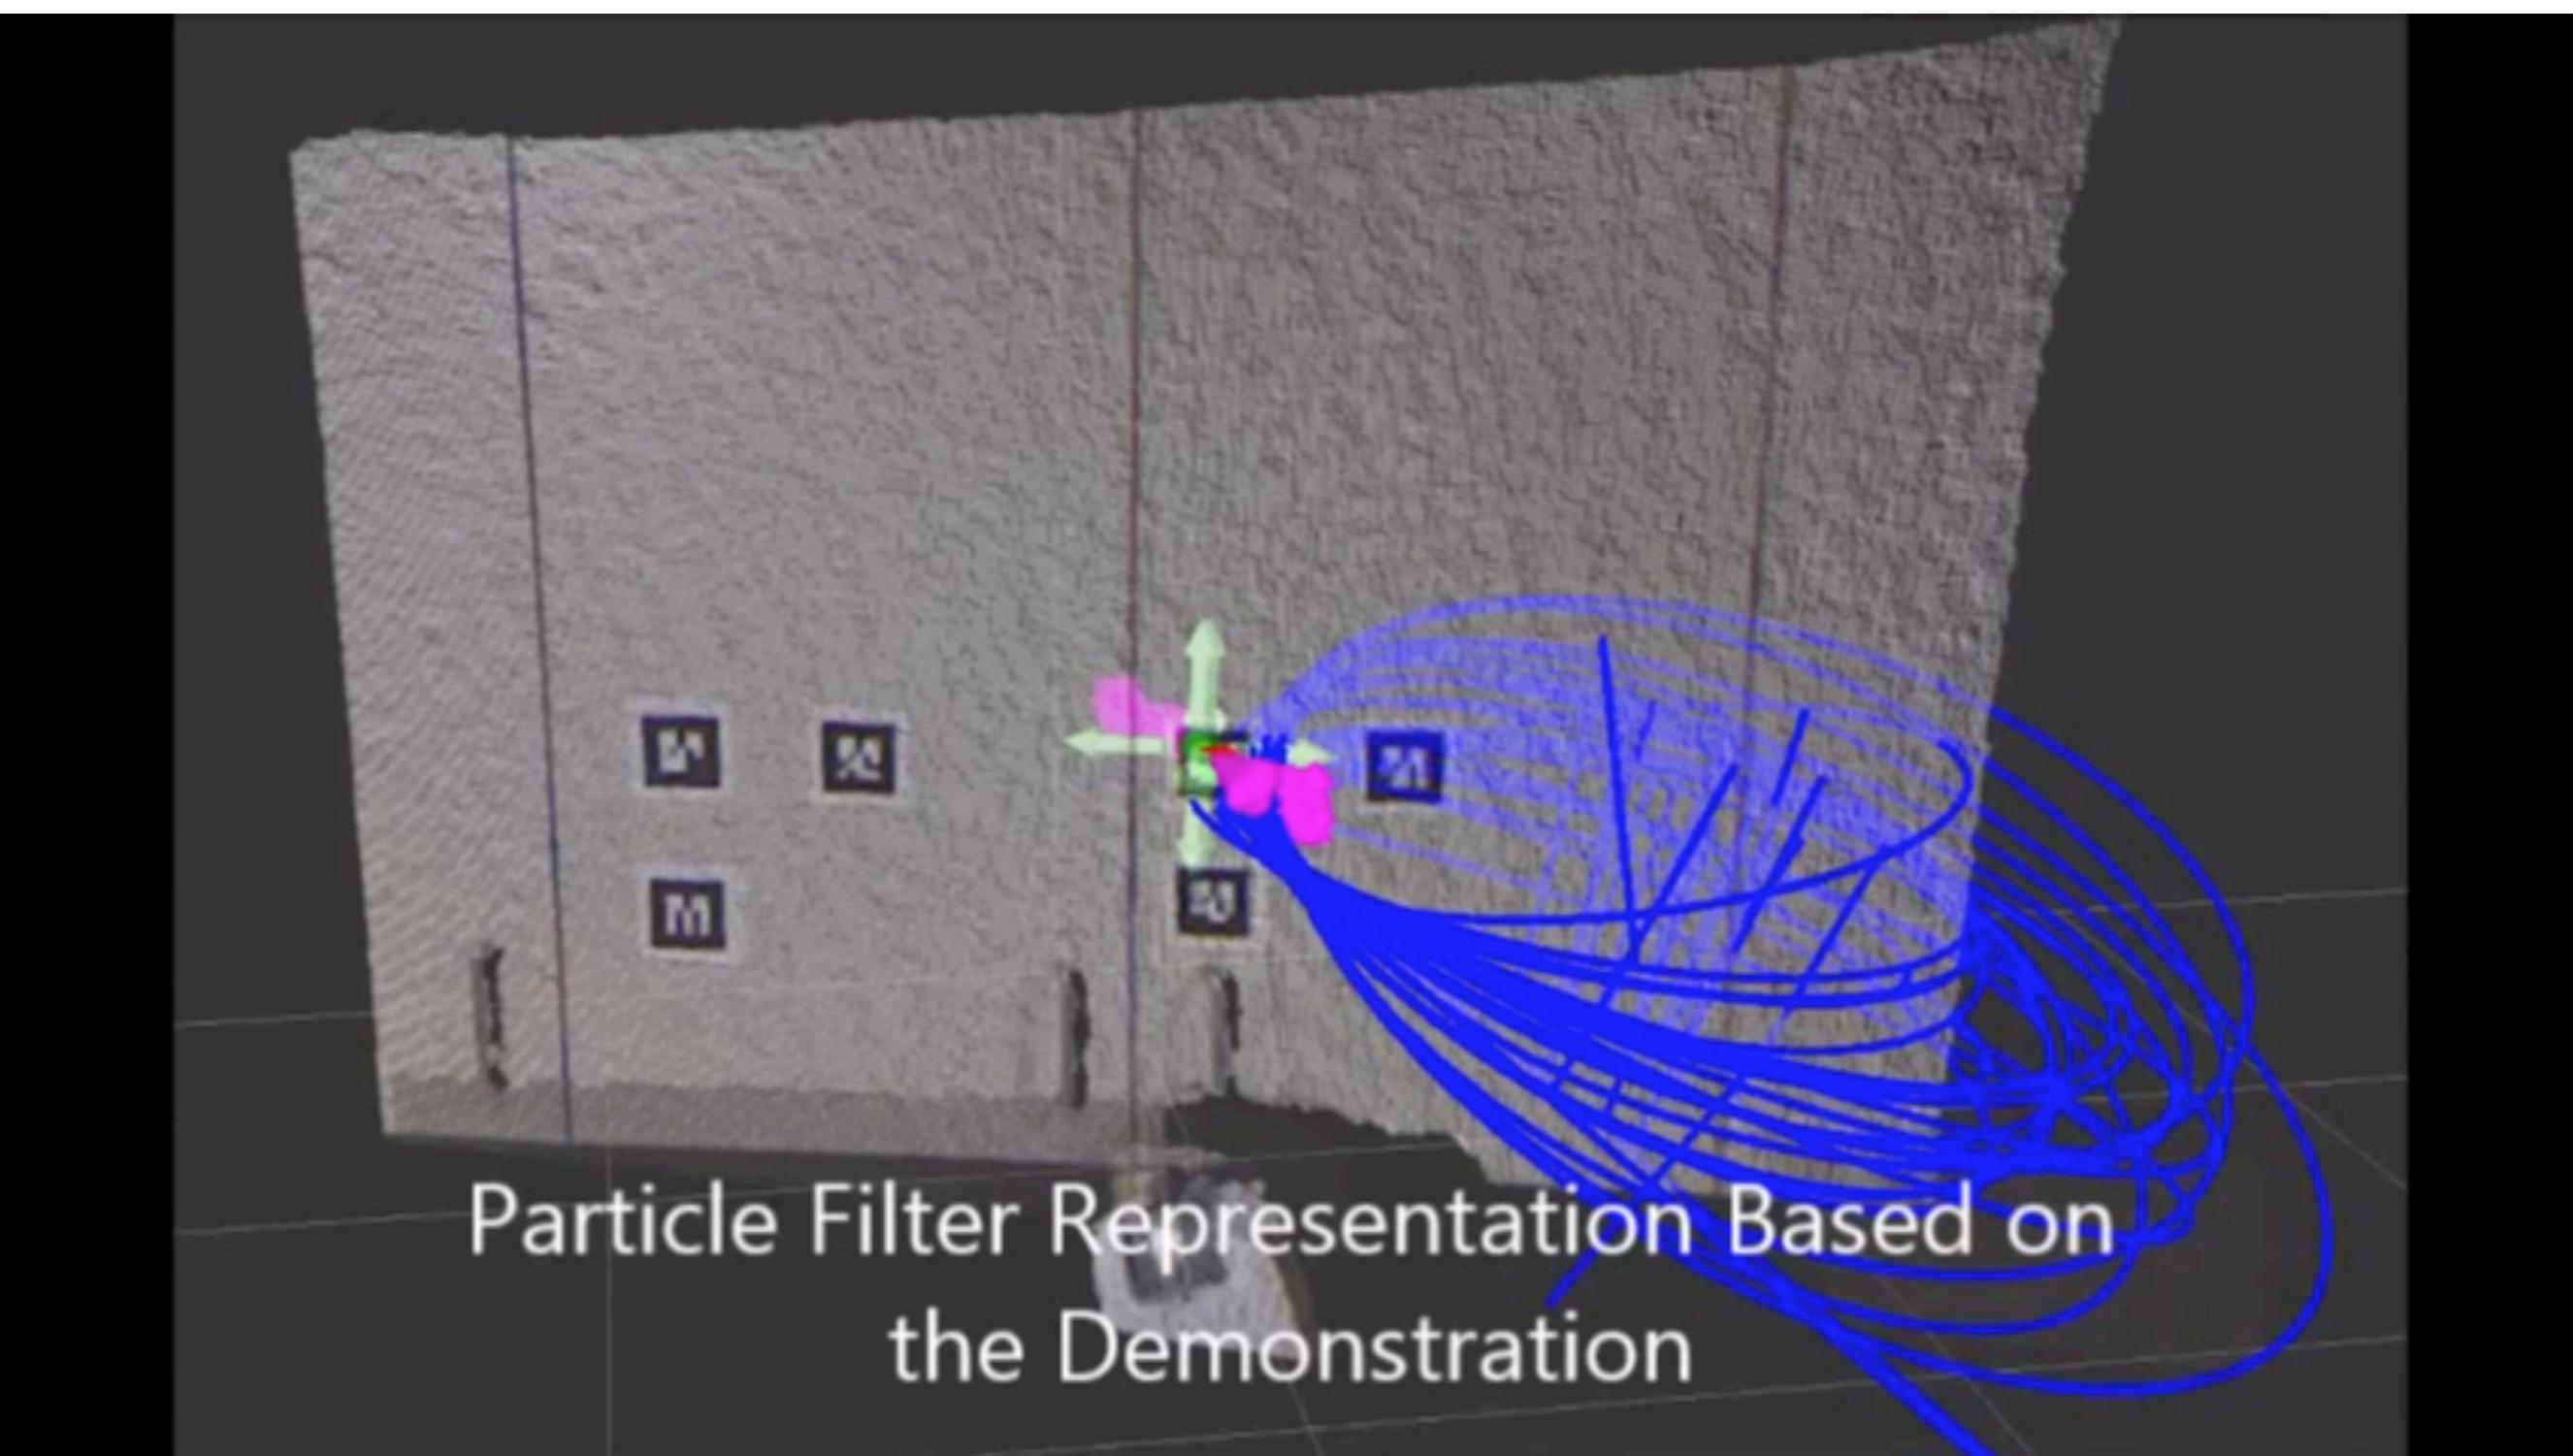
---

- Robot can perform **actions** to acquire **more information**
  - ▶ Change point of view to get a different perspective
  - ▶ Perform actions and observe effects - **interactive perception**
- **Interactive perception** more informative than passive
  - ▶ Additional sensor modalities, e.g., tactile
  - ▶ Infer latent properties of environment, e.g., mass
  - ▶ Overcome visual ambiguity
- **Additional information often basis for self-supervised learning**

# Segmentation Example



# Kinematics Example



Particle Filter Representation Based on  
the Demonstration

# Material Property Example

---



---

# **State Transition Models**

---

# Transition Models

- Transition probability

$$p(s_{t+1} | a_t, s_t)$$

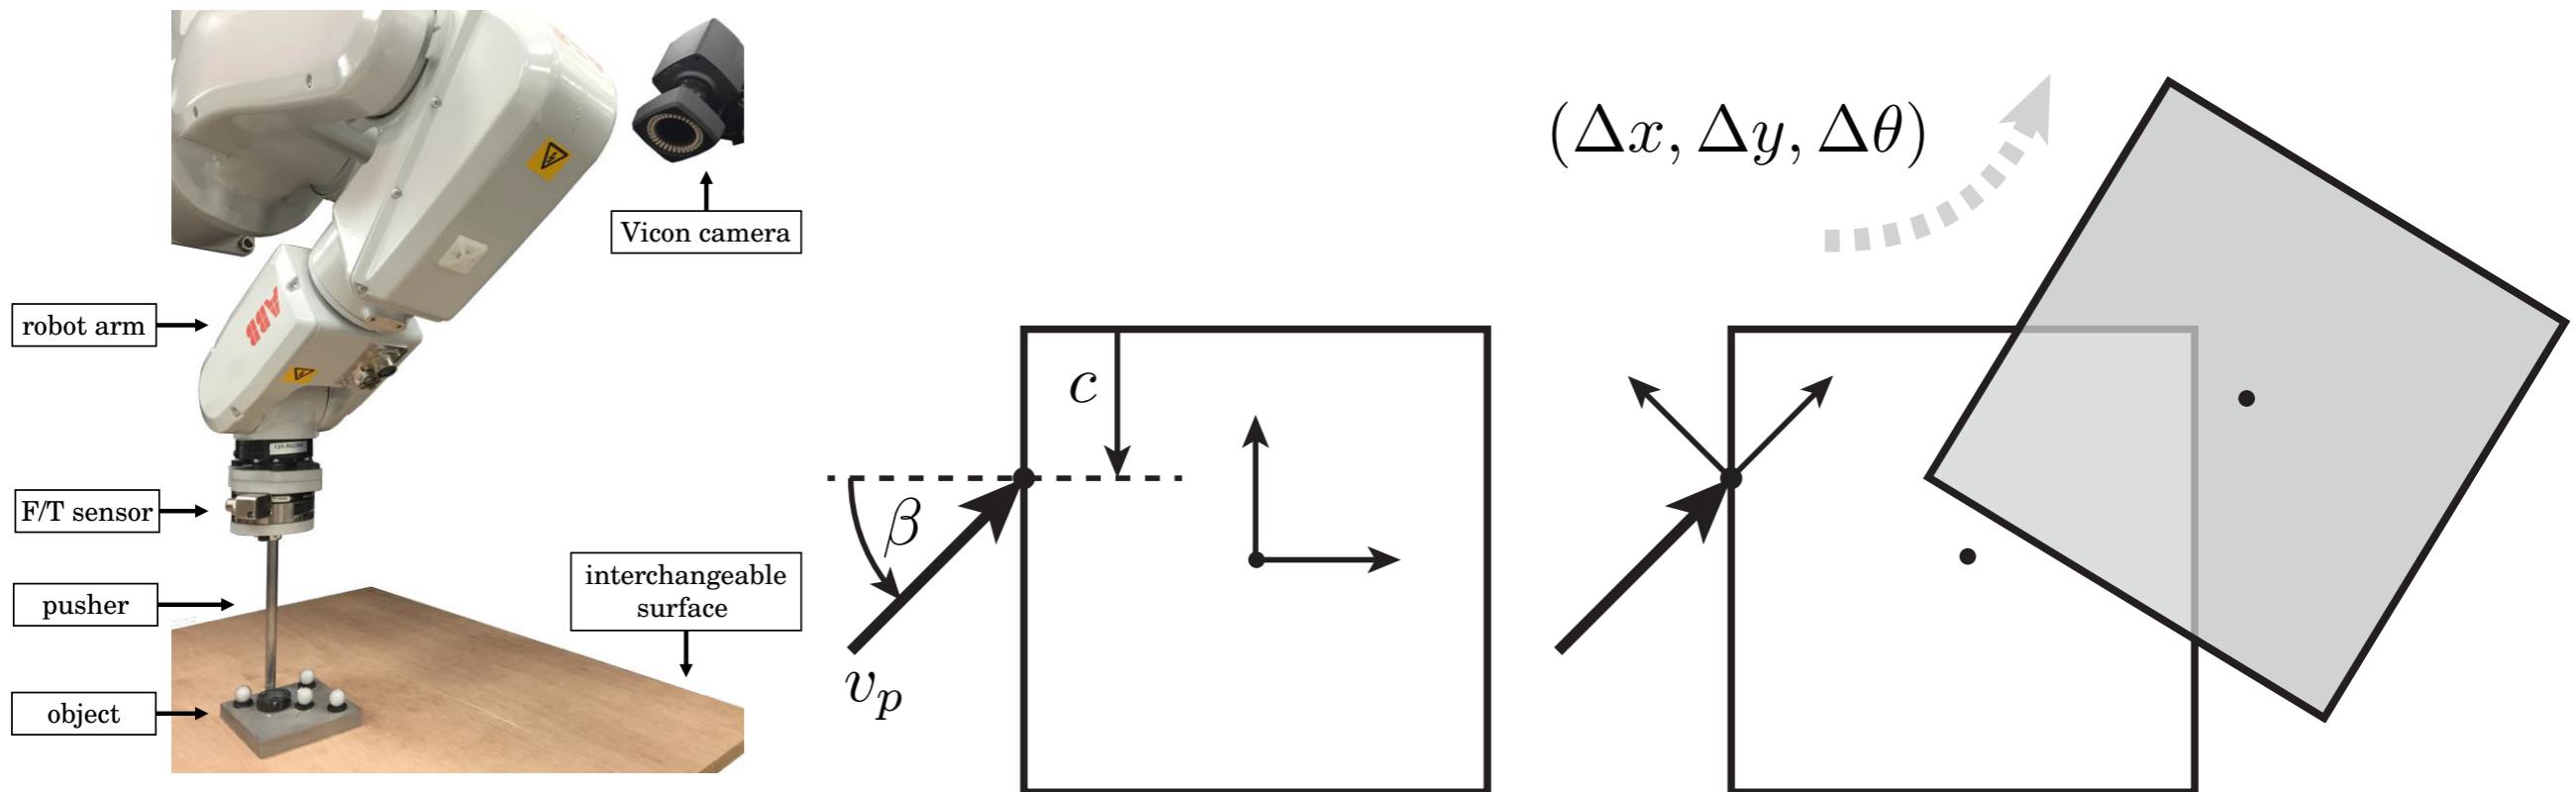
- Regression problem for continuous states



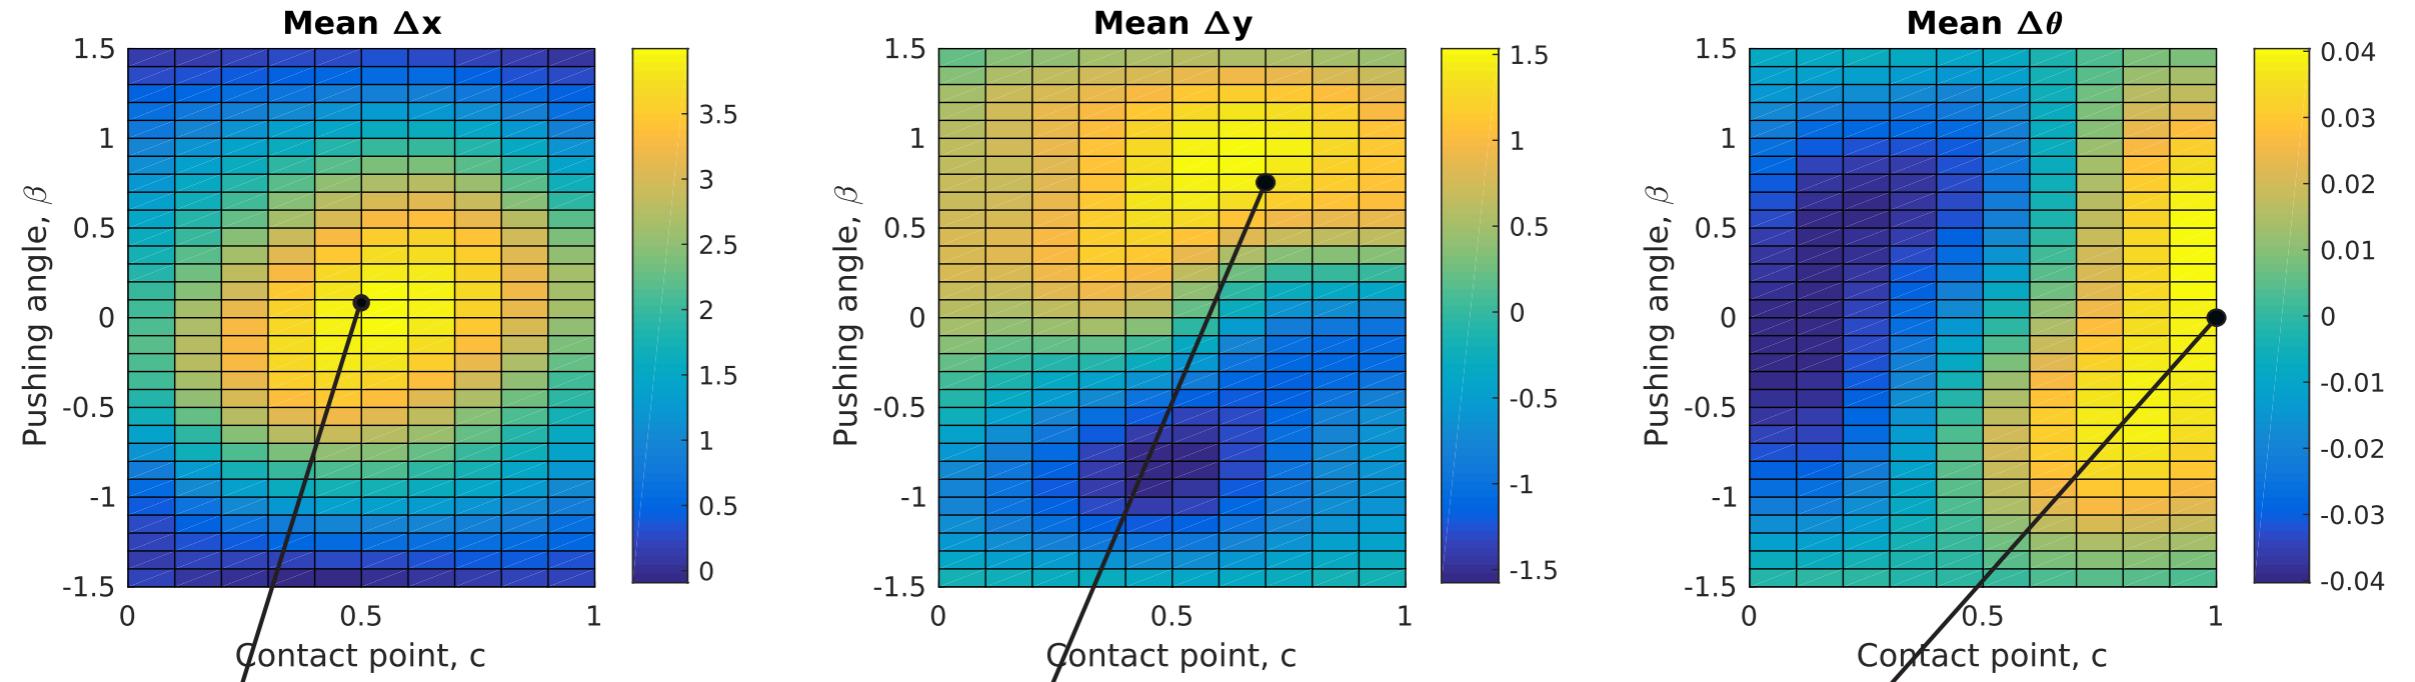
- Robot can acquire these training samples autonomously
  - ▶ Observe state, perform an action, and observe the next state
  - ▶ Samples can be acquired independent of a specific task
- Capture transitions using various different models
  - ▶ Linear regression, Gaussian process regression, neural networks,...

# Model for Pushing

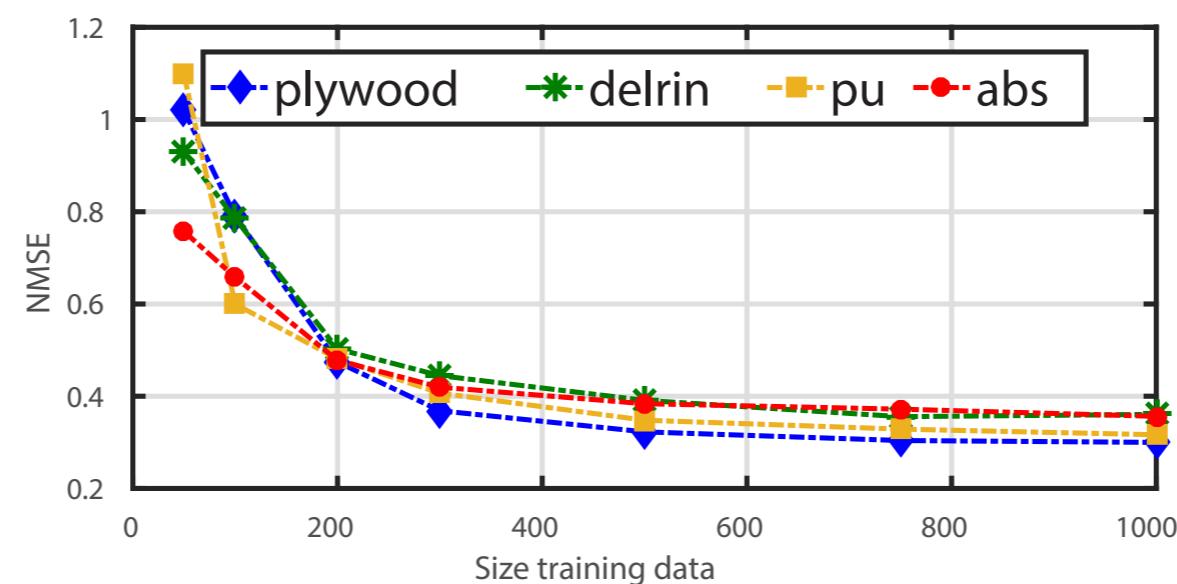
- Use a **Gaussian Process** to model planar pushing effects
  - ▶ **Input:** Contact point position, angle, and velocity
  - ▶ **Output:** contact-frame-relative shifts in position and orientation



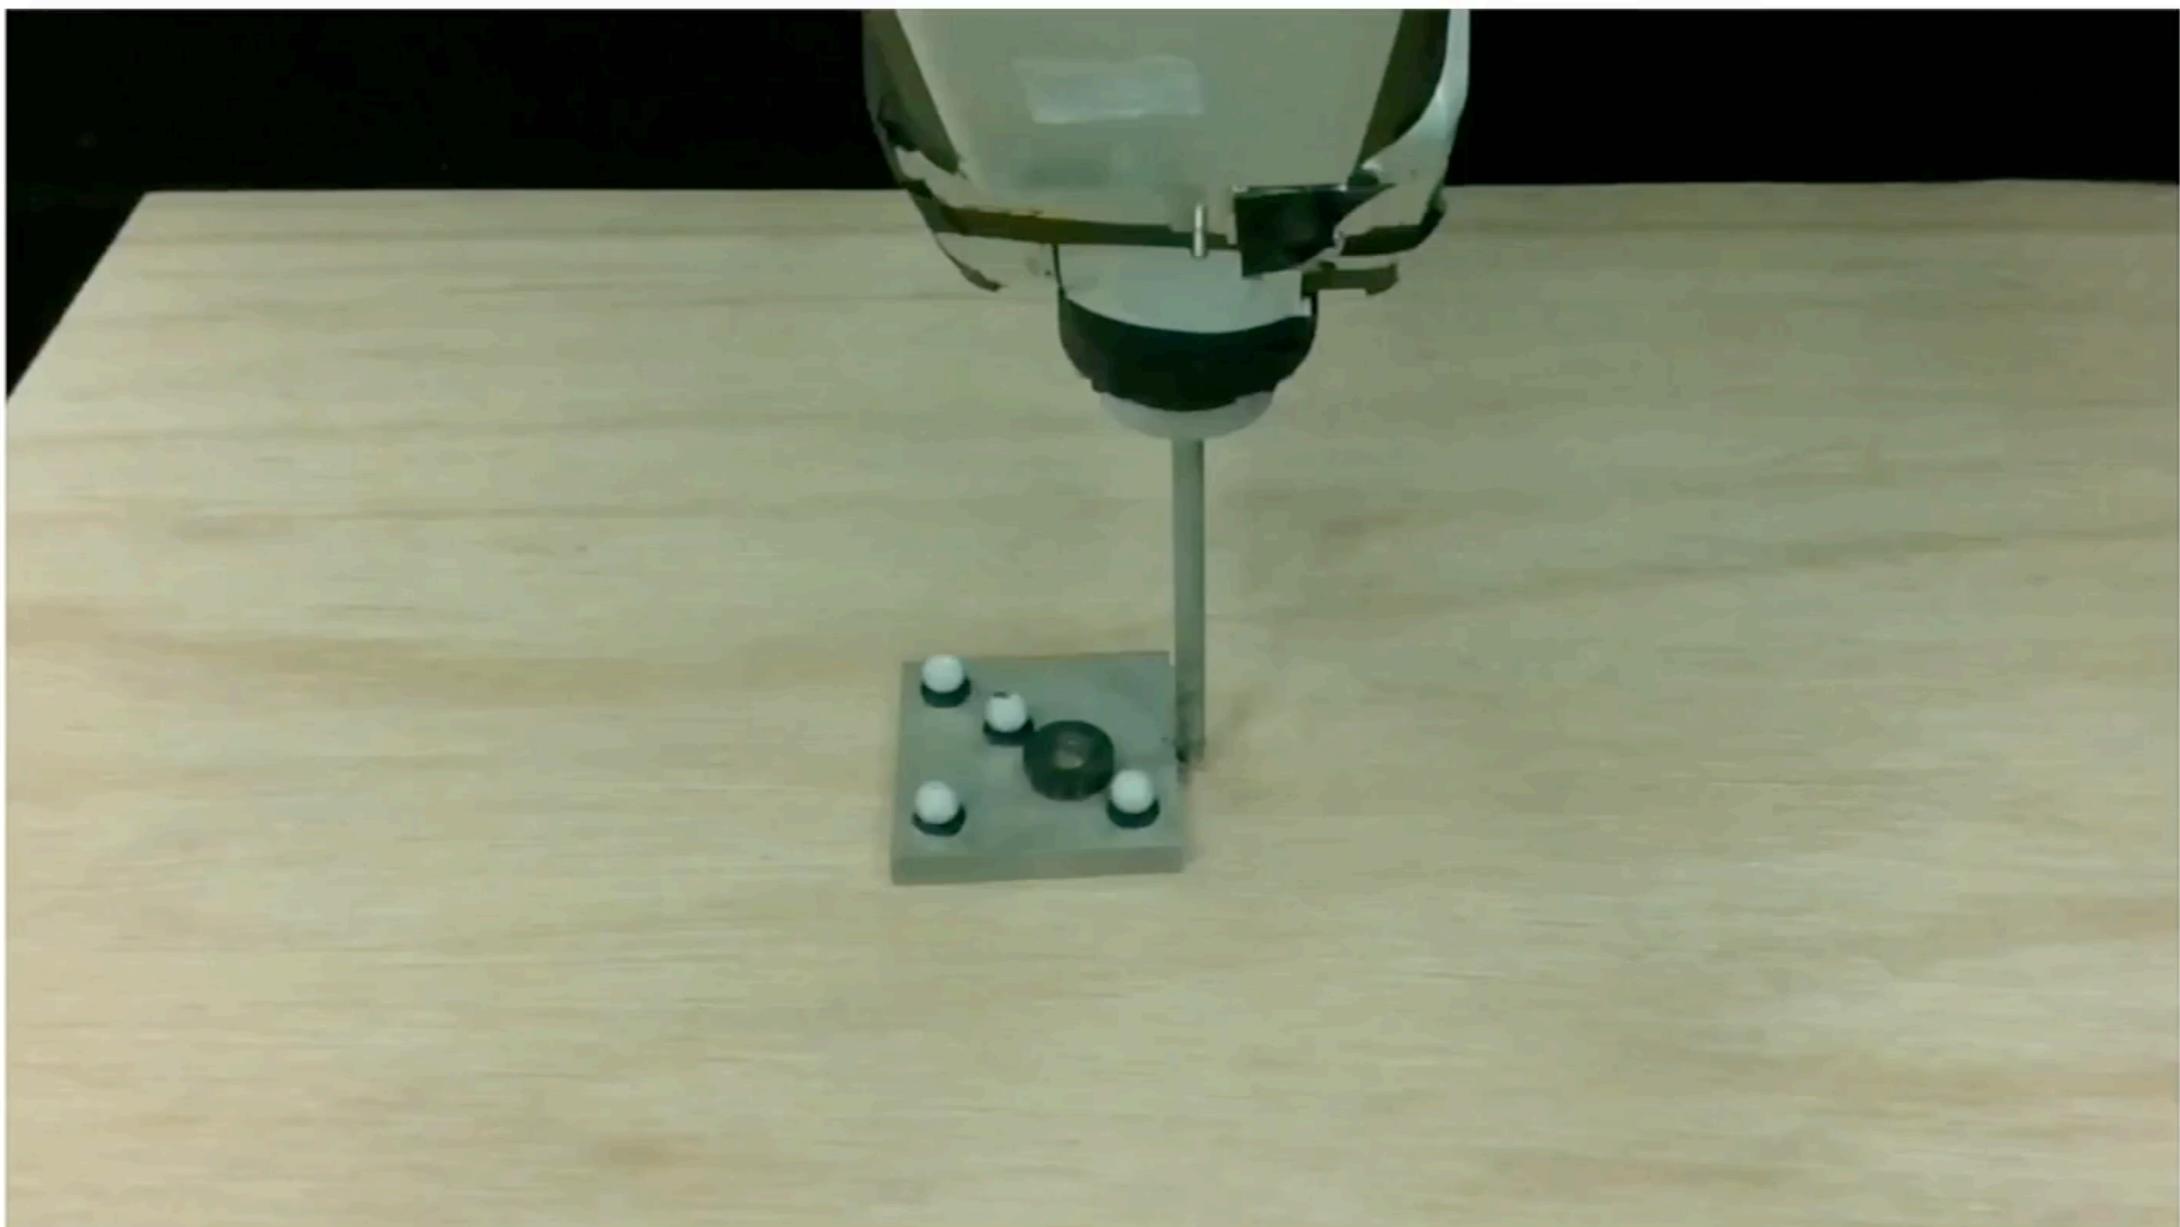
# Model for Pushing



$$\text{NMSE} = \frac{\sum_{j=1}^m (y_j - \hat{y}_j)^2}{\sum_{j=1}^m (y_j - \bar{y})^2}$$

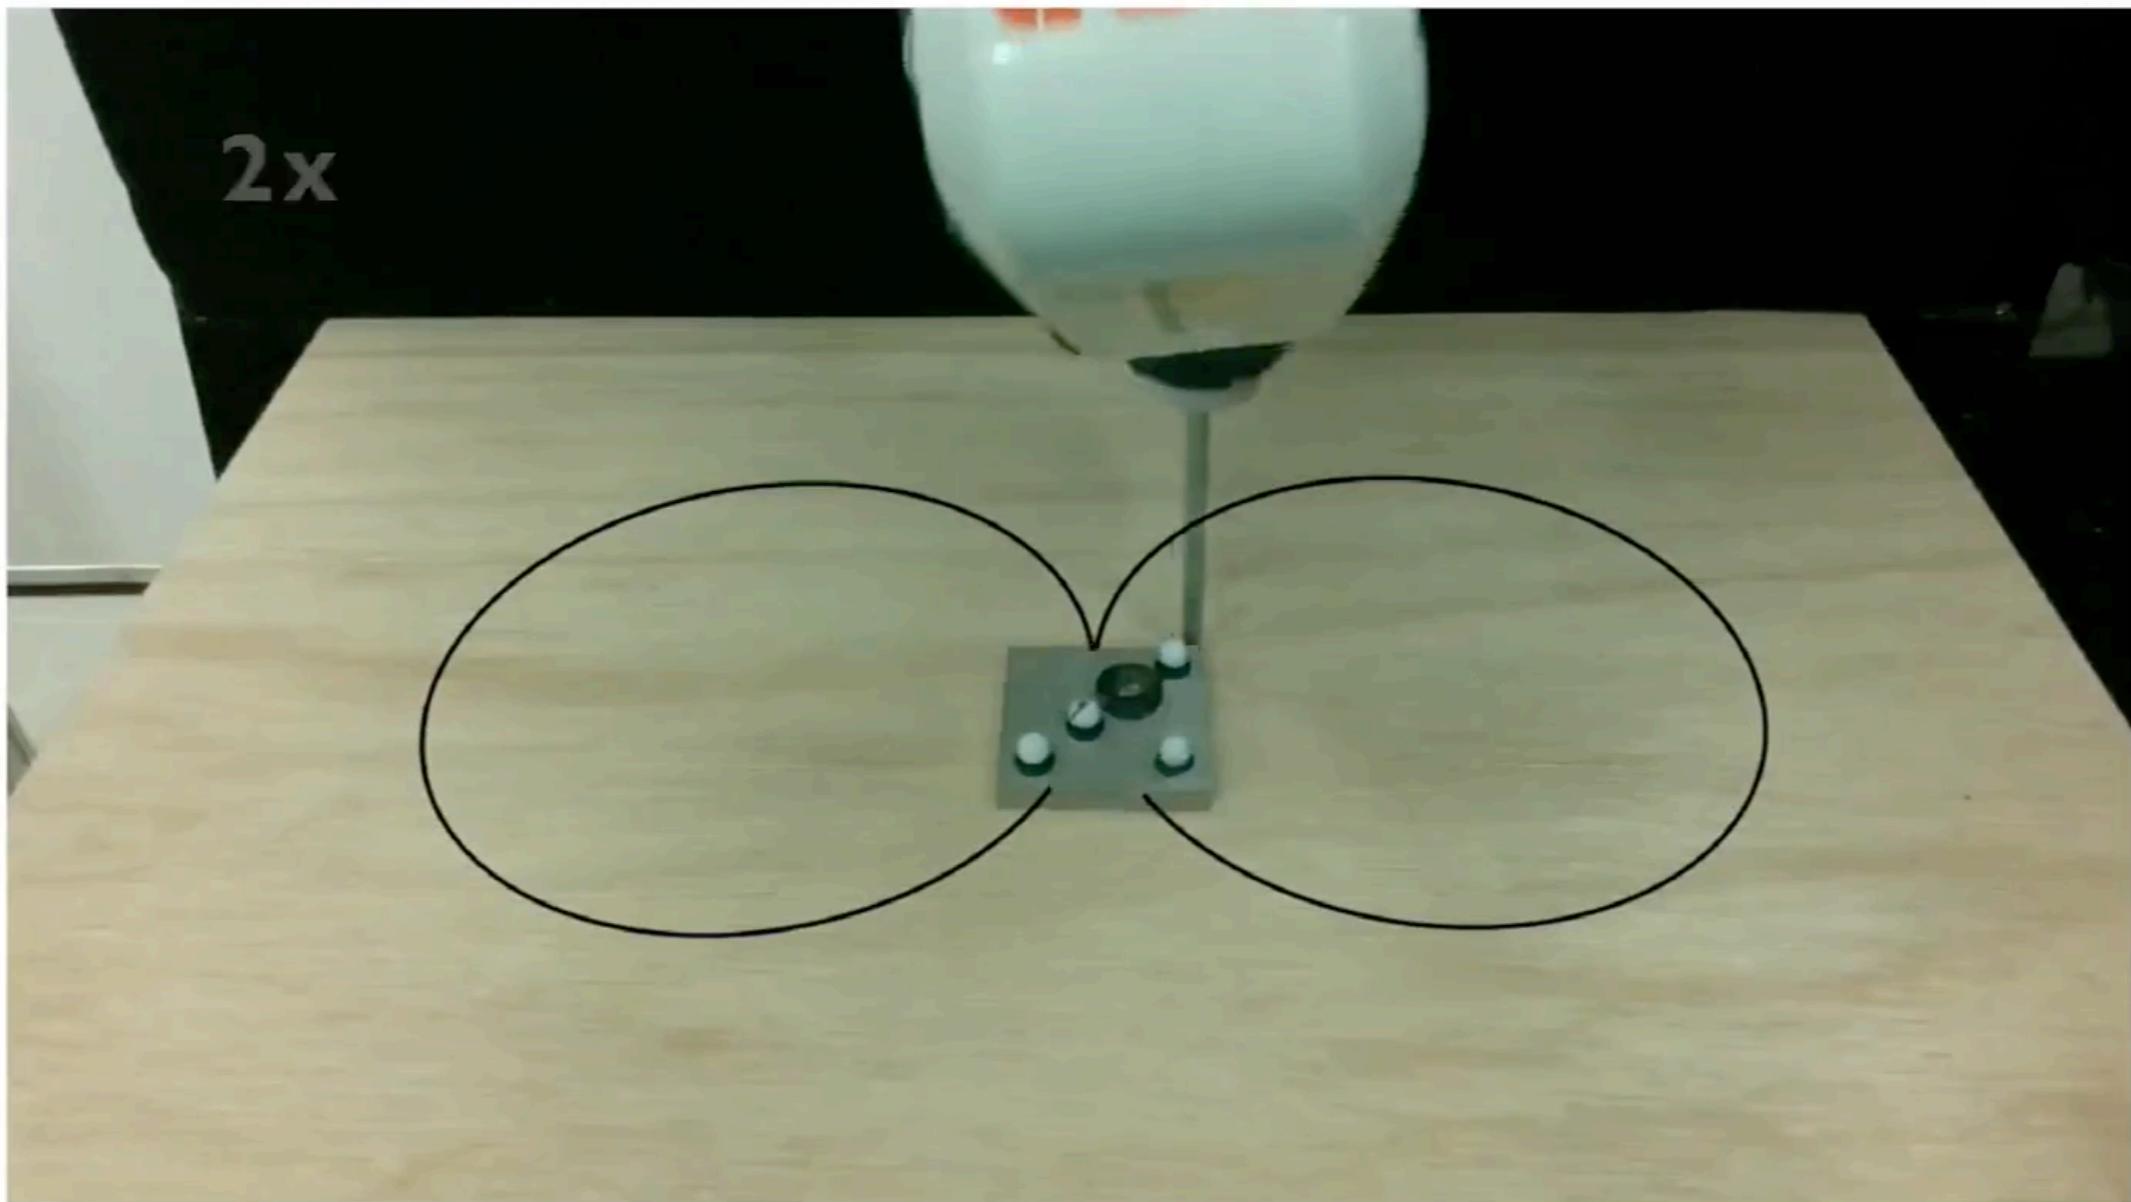


## Data collection



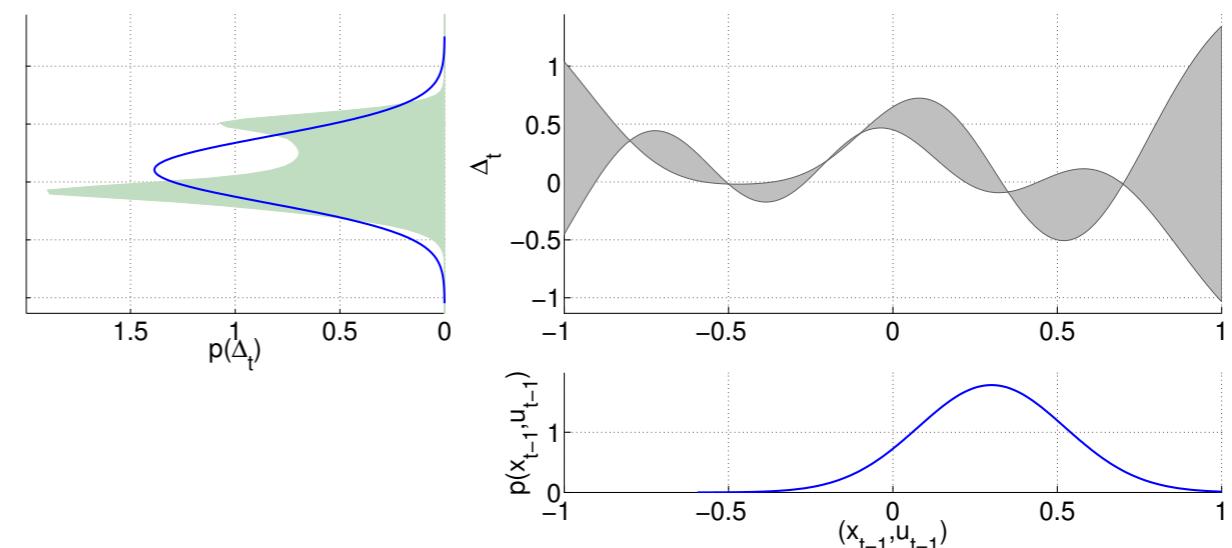
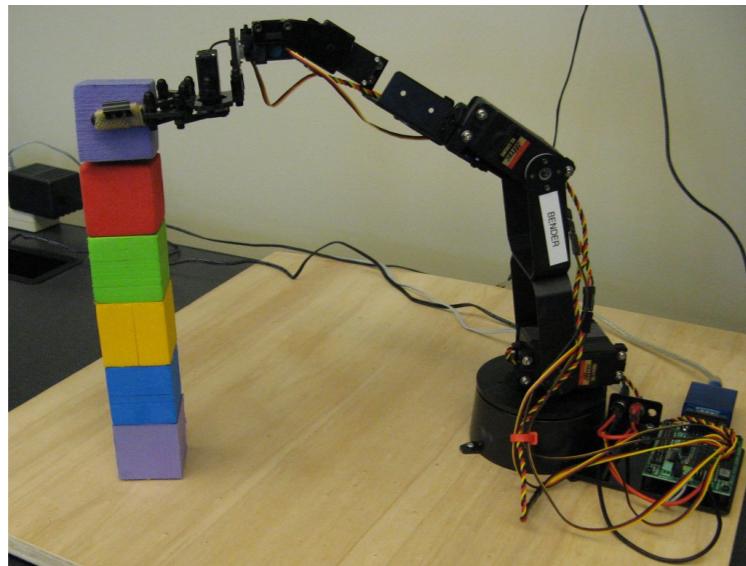
# Model for Pushing

## A Data-Efficient Approach to Precise and Controlled Pushing



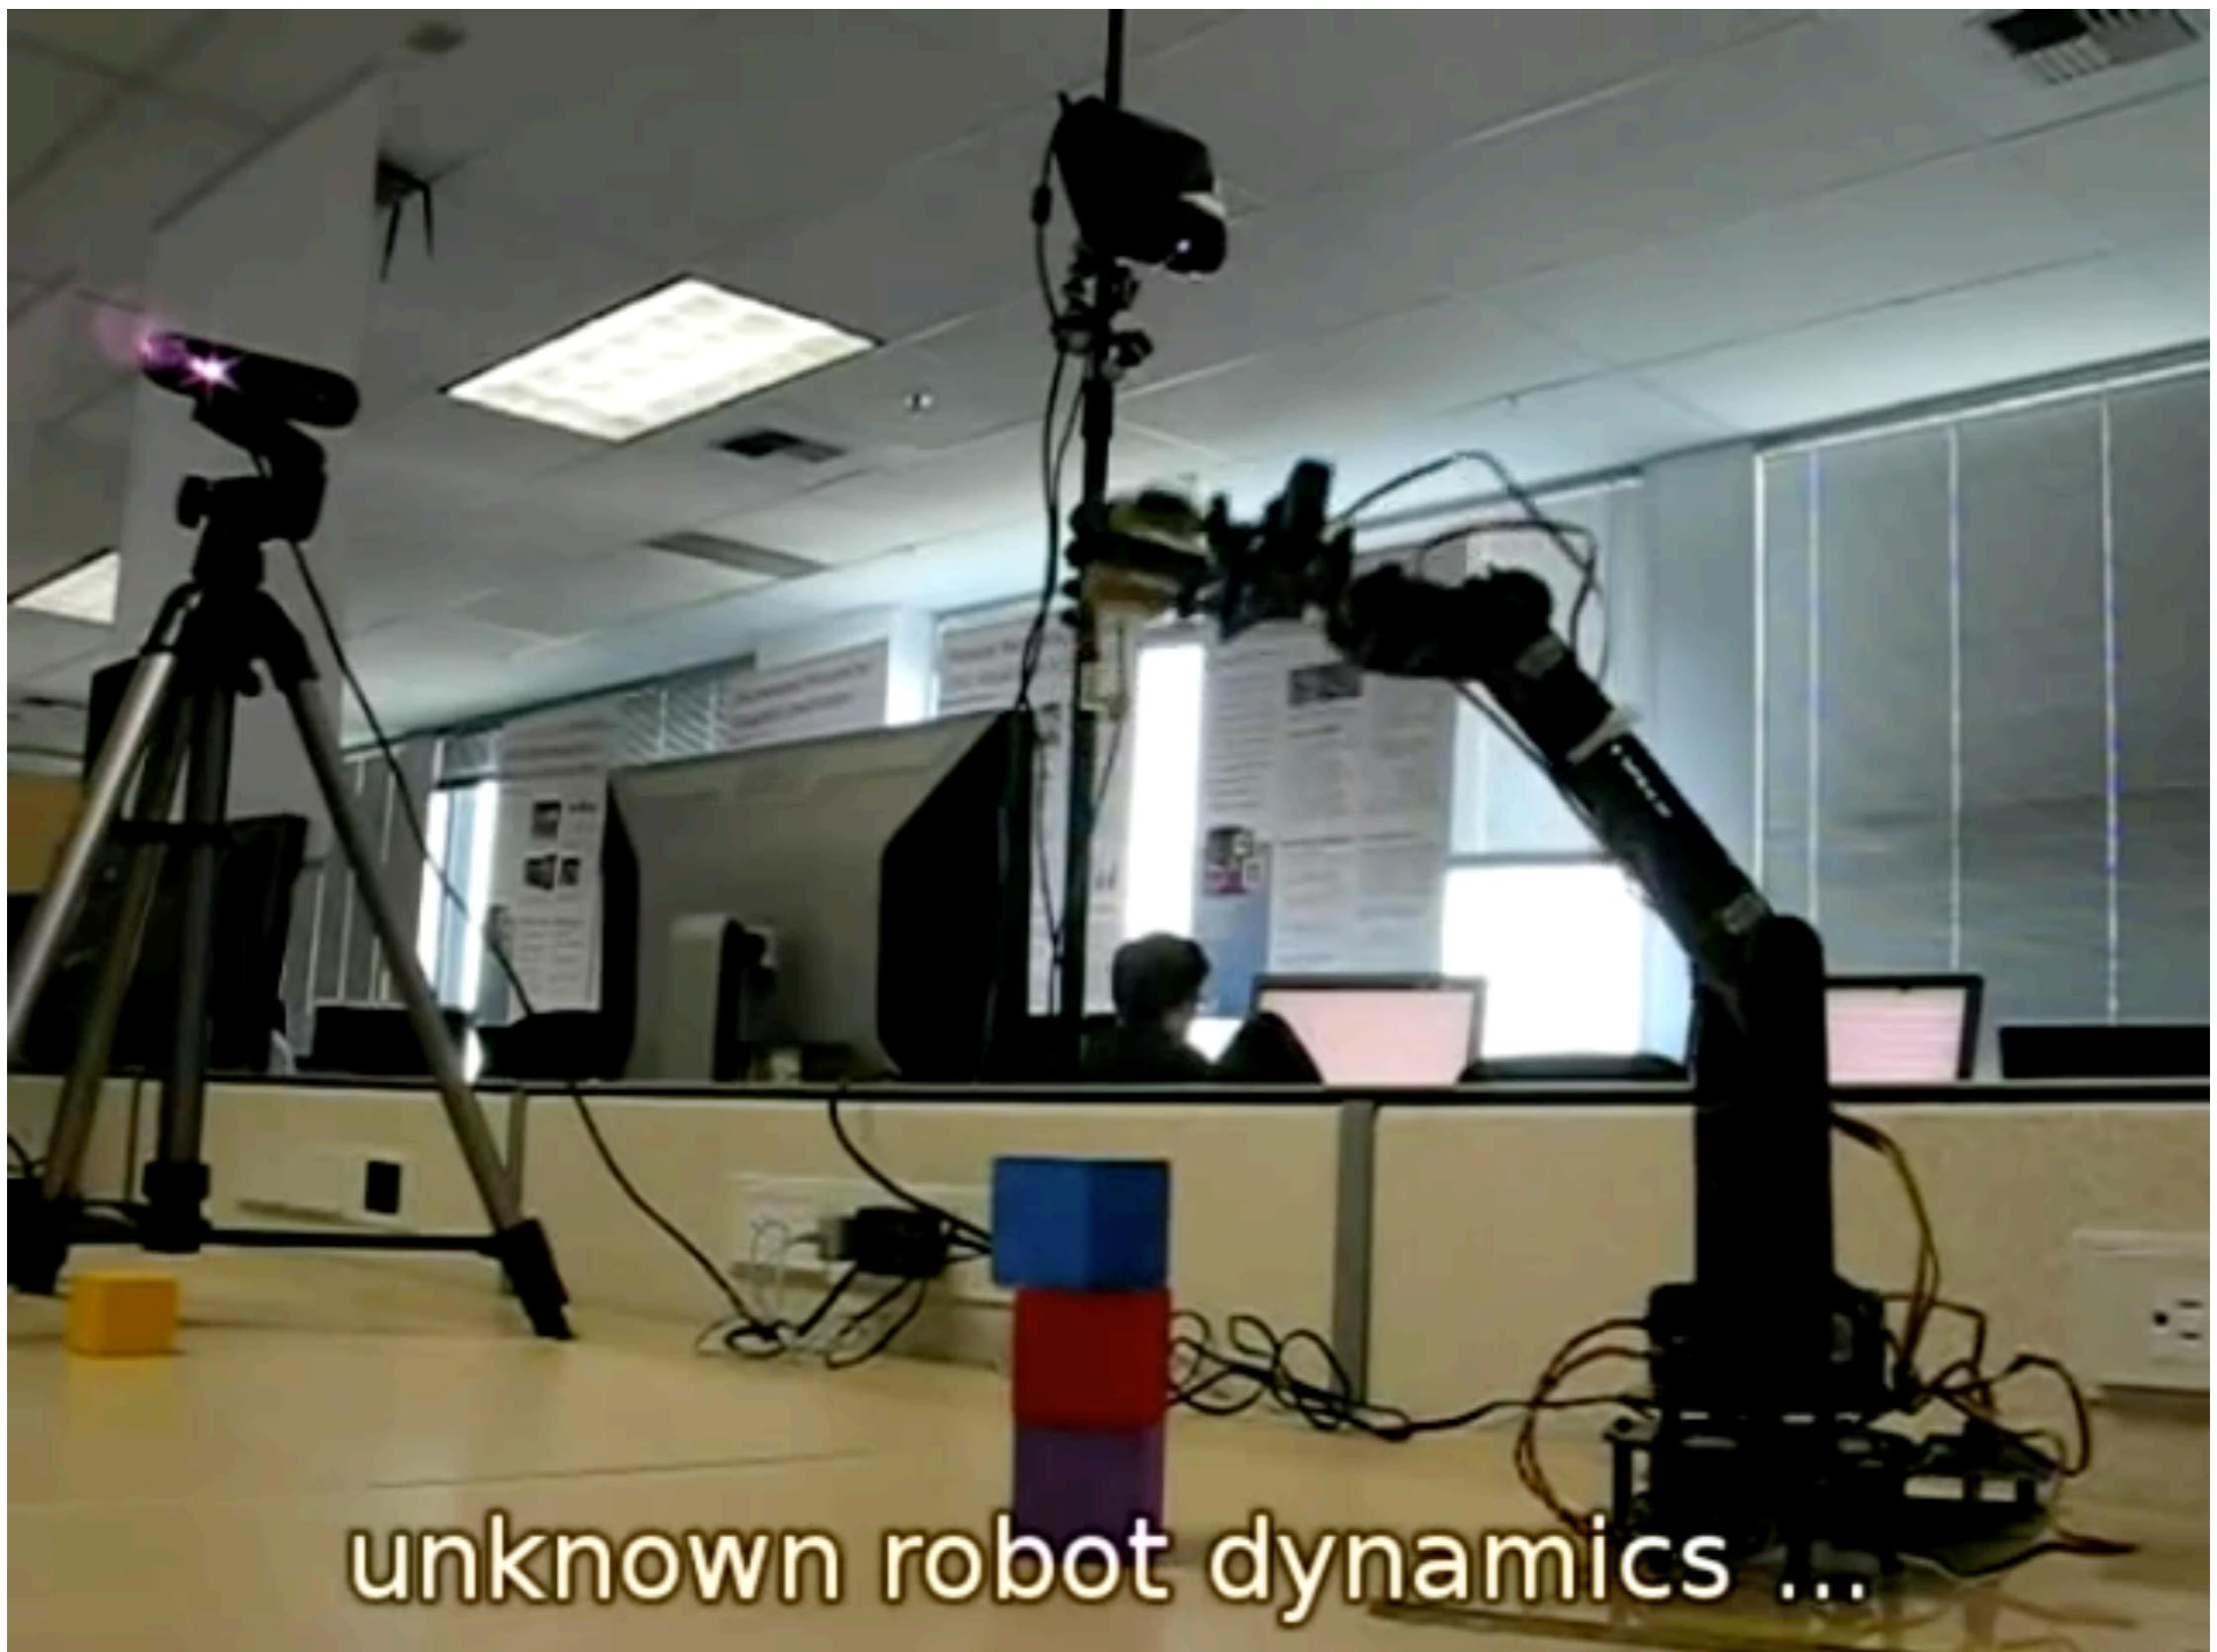
# Vision-based Block Stacking Example

- Learn (local) models to stack blocks
  - ▶ State (3D): centroid of block using blob tracking
  - ▶ Action (4D): pulse widths for first four joint servos



- Models can capture stochasticity/uncertainty
- Use model to compute an optimal policy

# Vision-based Block Stacking Example



---

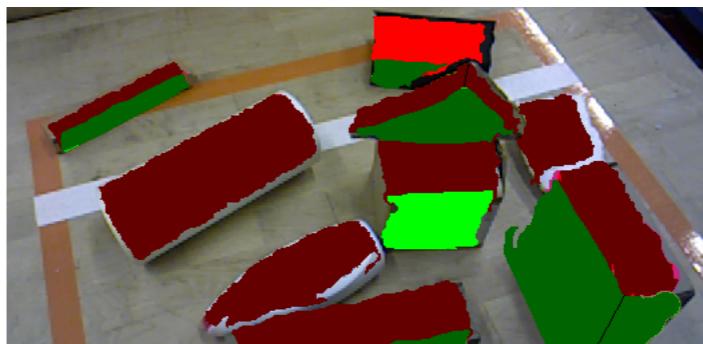
# Affordances

---

# Affordances

- Objects **afford** actions/manipulations to an **agent**
  - ▶ Apple affords grasping, picking, eating, throwing... to a human
  - ▶ Apple affords grasping, picking, throwing... to a robot

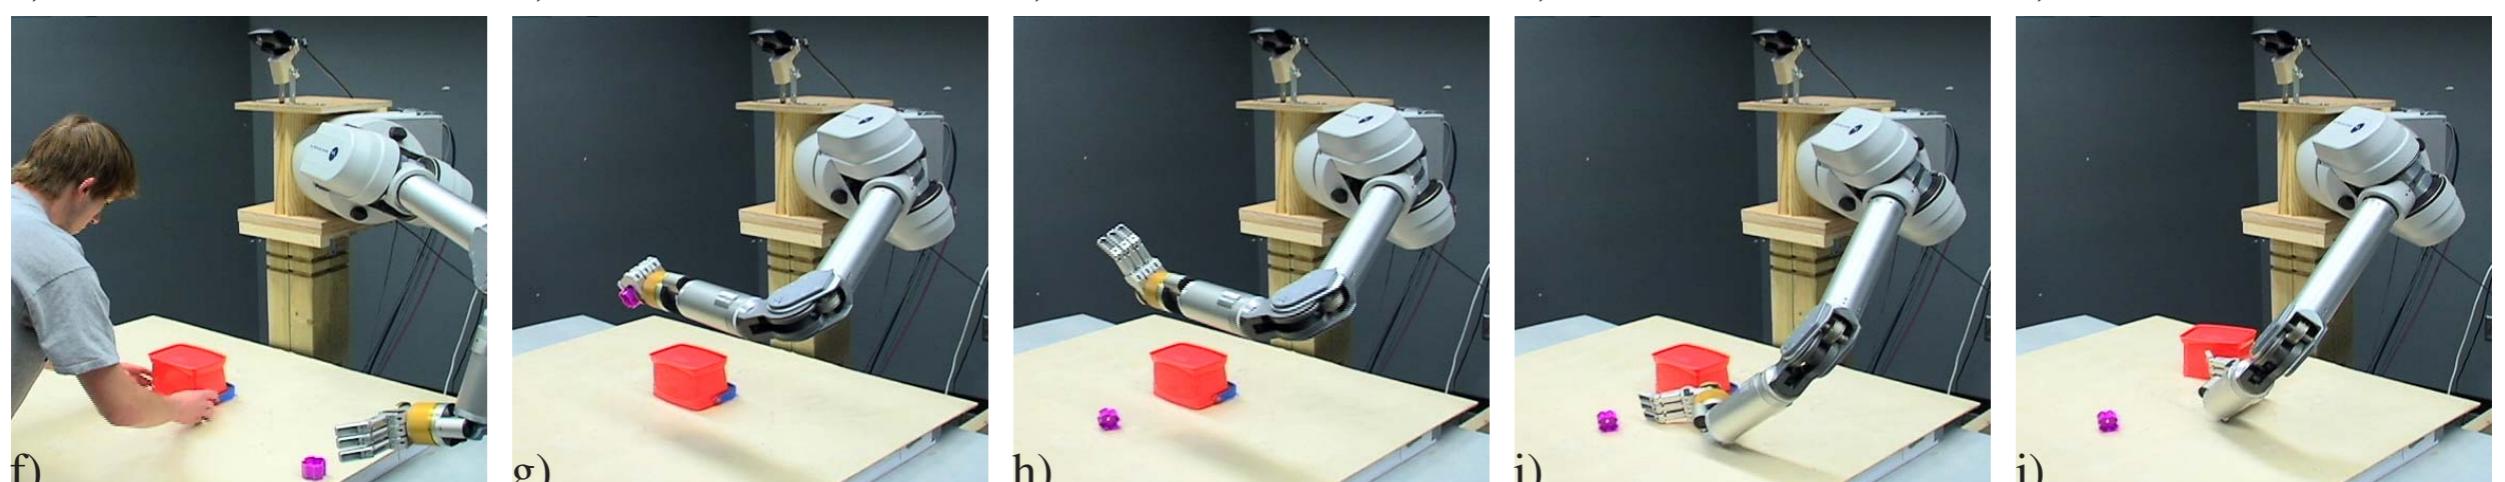
**Surfaces that afford pushing  
to slide by a manipulator**



- Affordances connect **objects** and **agents** to **action possibilities**
  - ▶ **Ground labels** (graspable, liftable, etc.) in robot's actions
  - ▶ **Robot can determine affordance label by performing action**

# Container Example

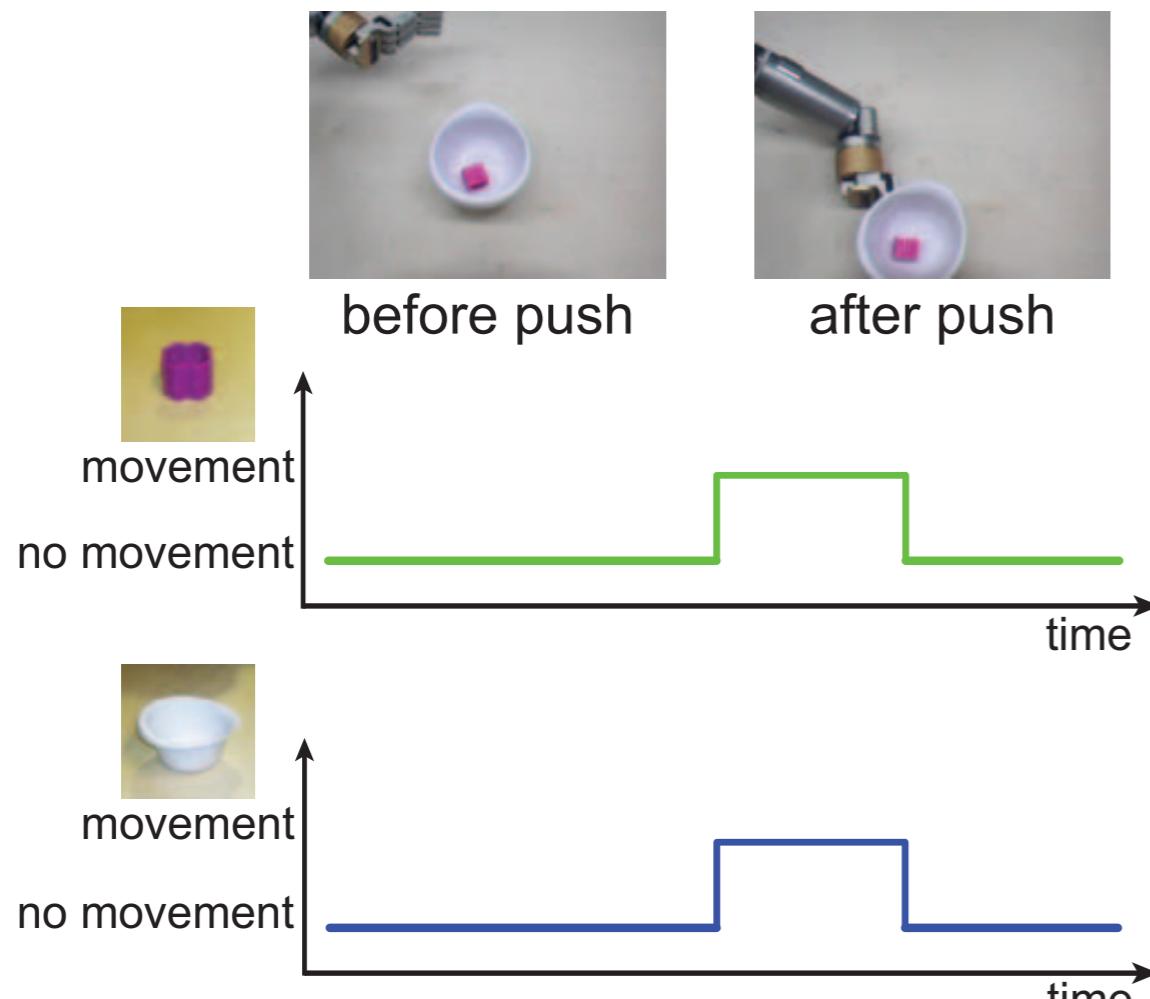
learn from interactions with containers & non-containers



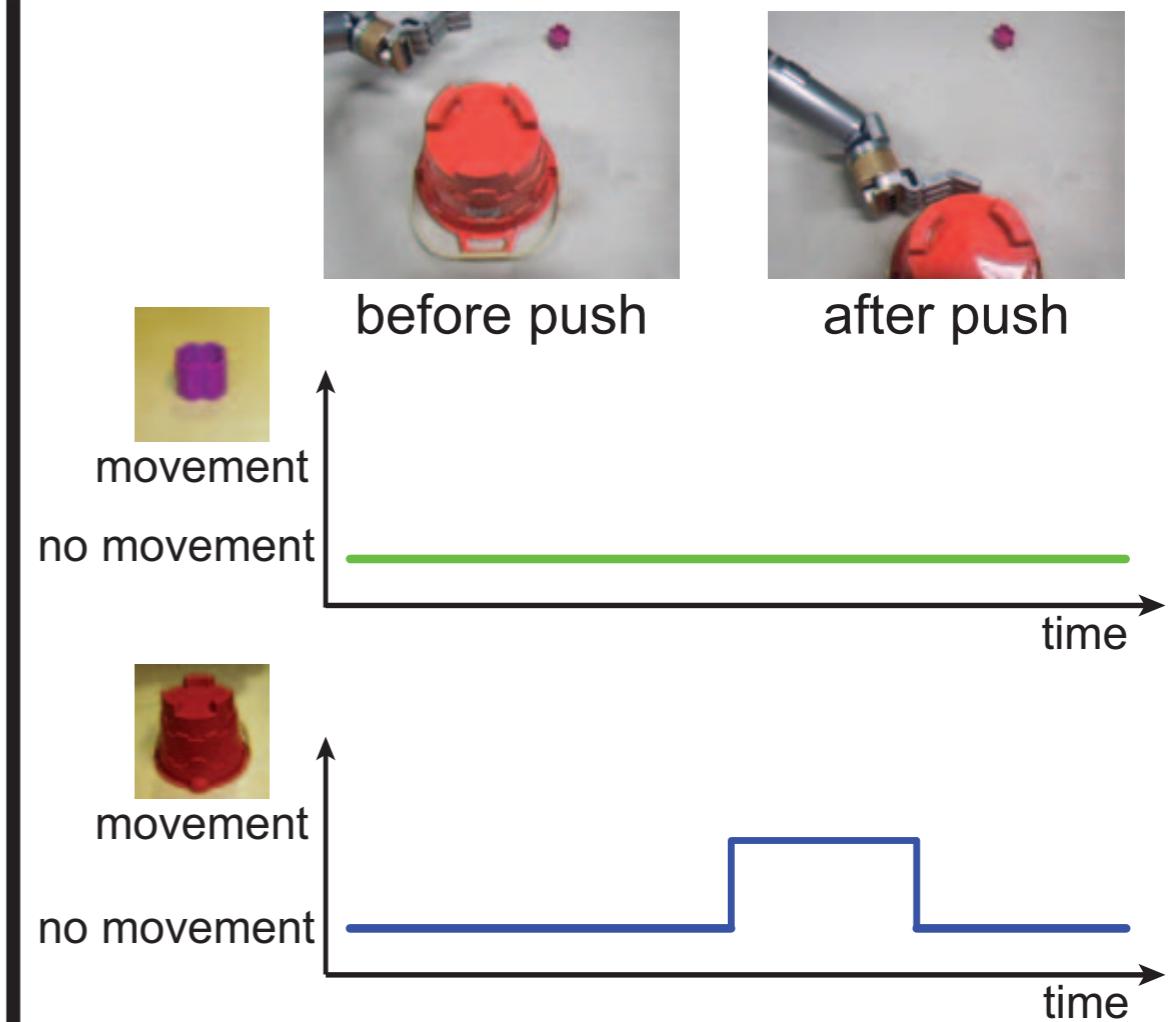
# Container Example

- Distinct classes of effects based on co-movement

Co-Movement



Separate Movement

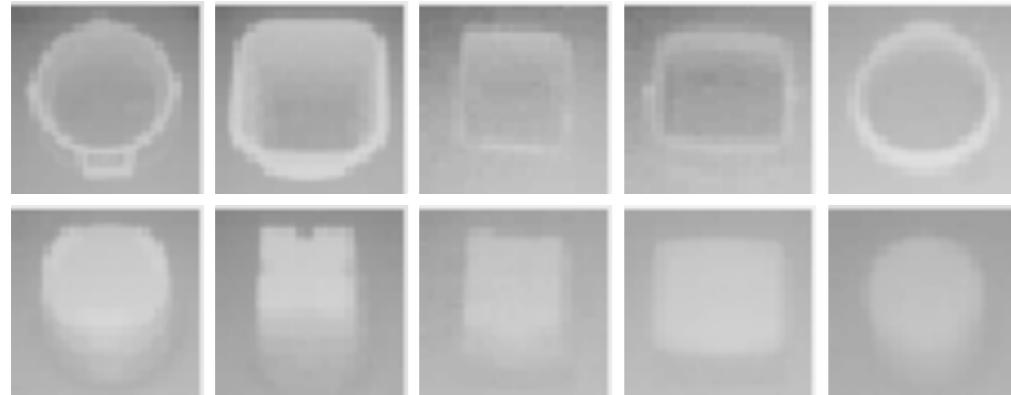


(a smaller third class corresponds to noisy tracking)

# Container Example

- Learn classifier for affordance (container or not)

Example Classifier Inputs



Classification Results

Novel Non-containers

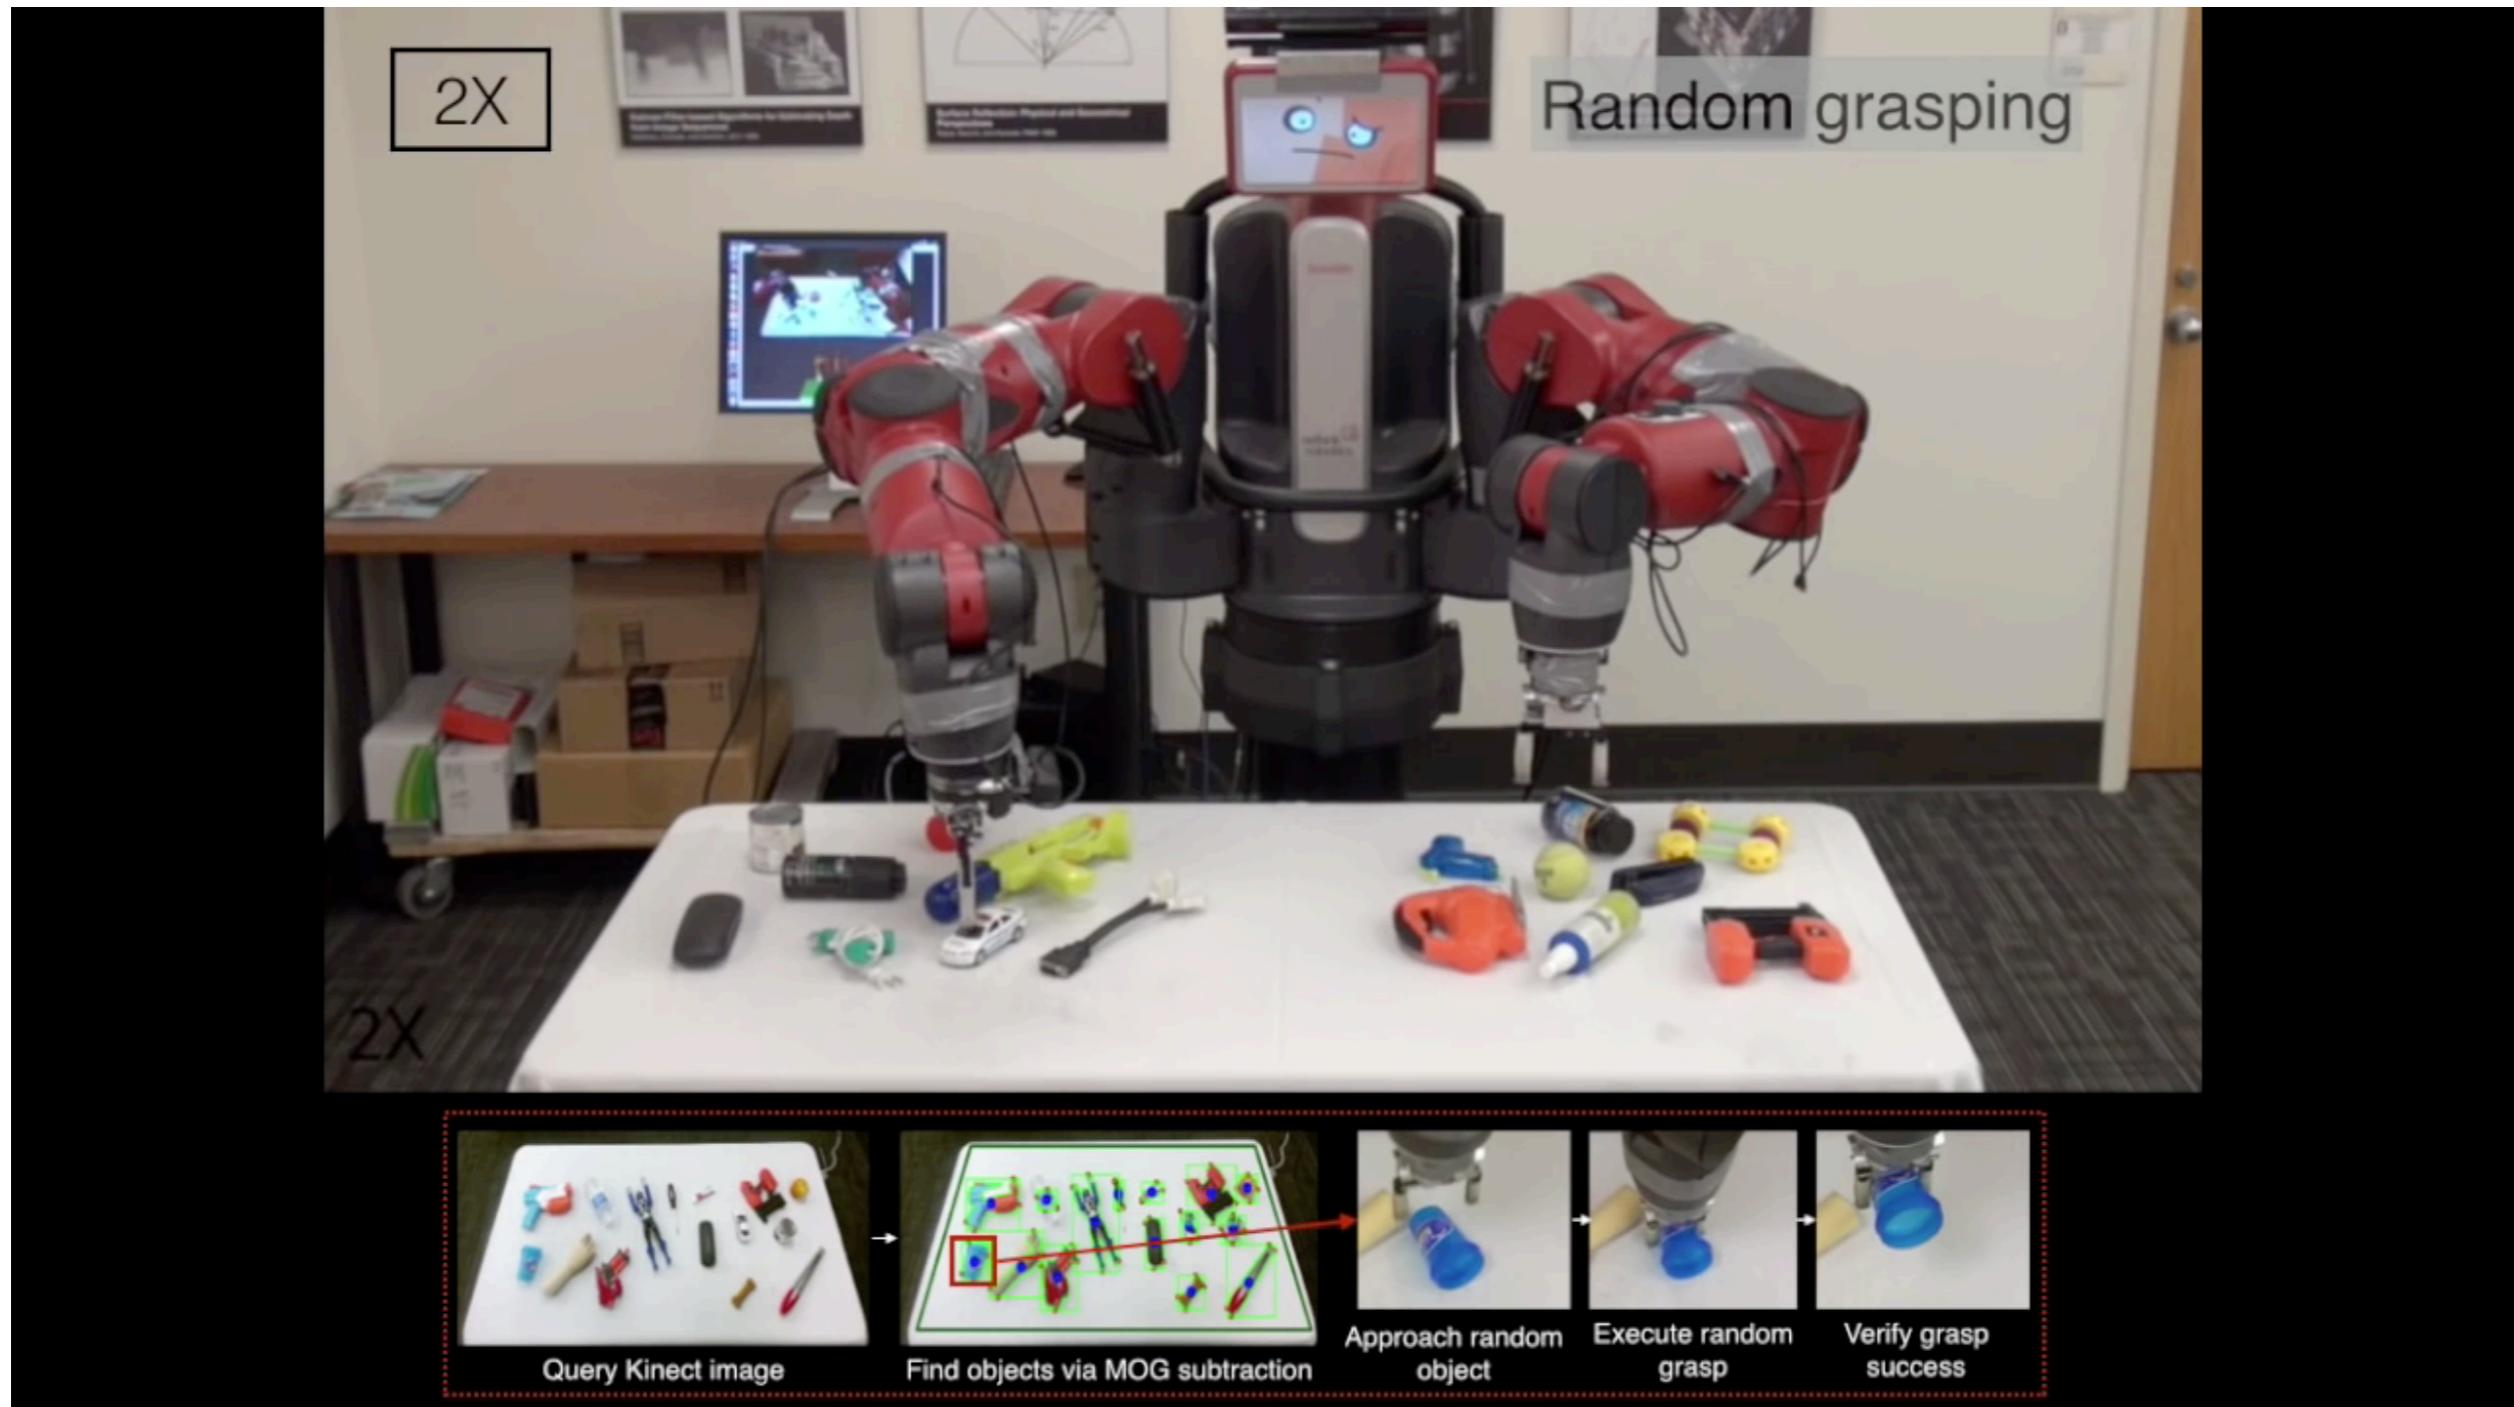


Novel Containers



Can learn **container vs. non-container**. What about **cup vs. mug**?

# Self-supervised Grasp Affordances

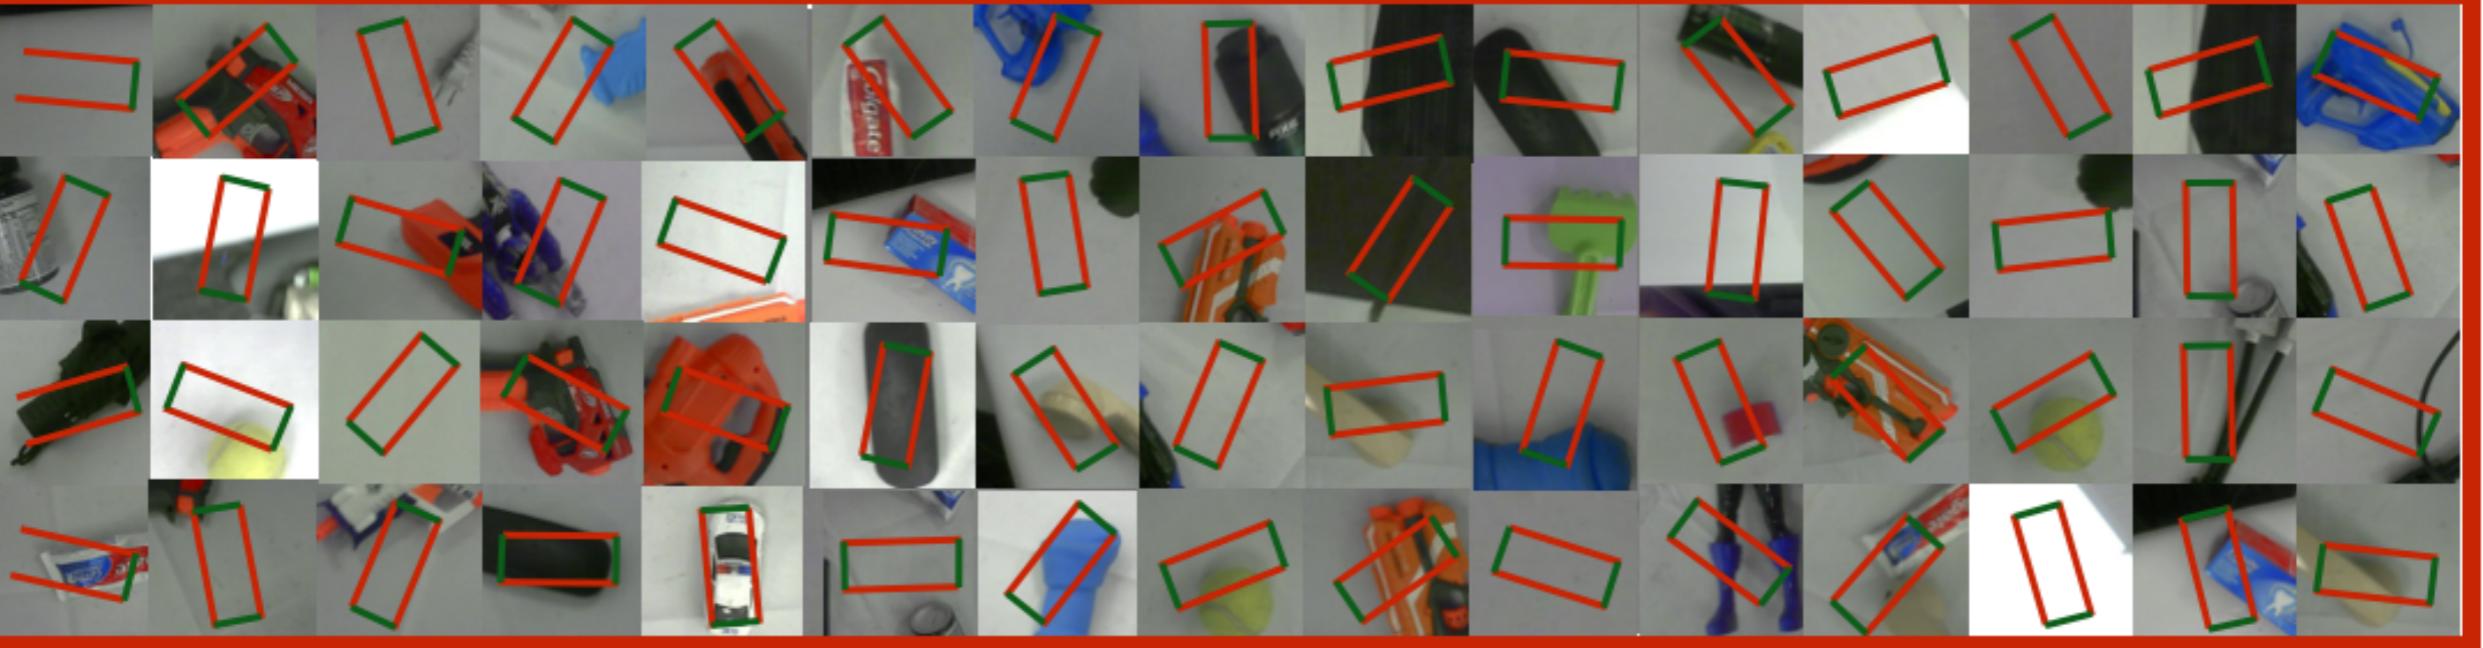


# Self-supervised Grasp Affordances

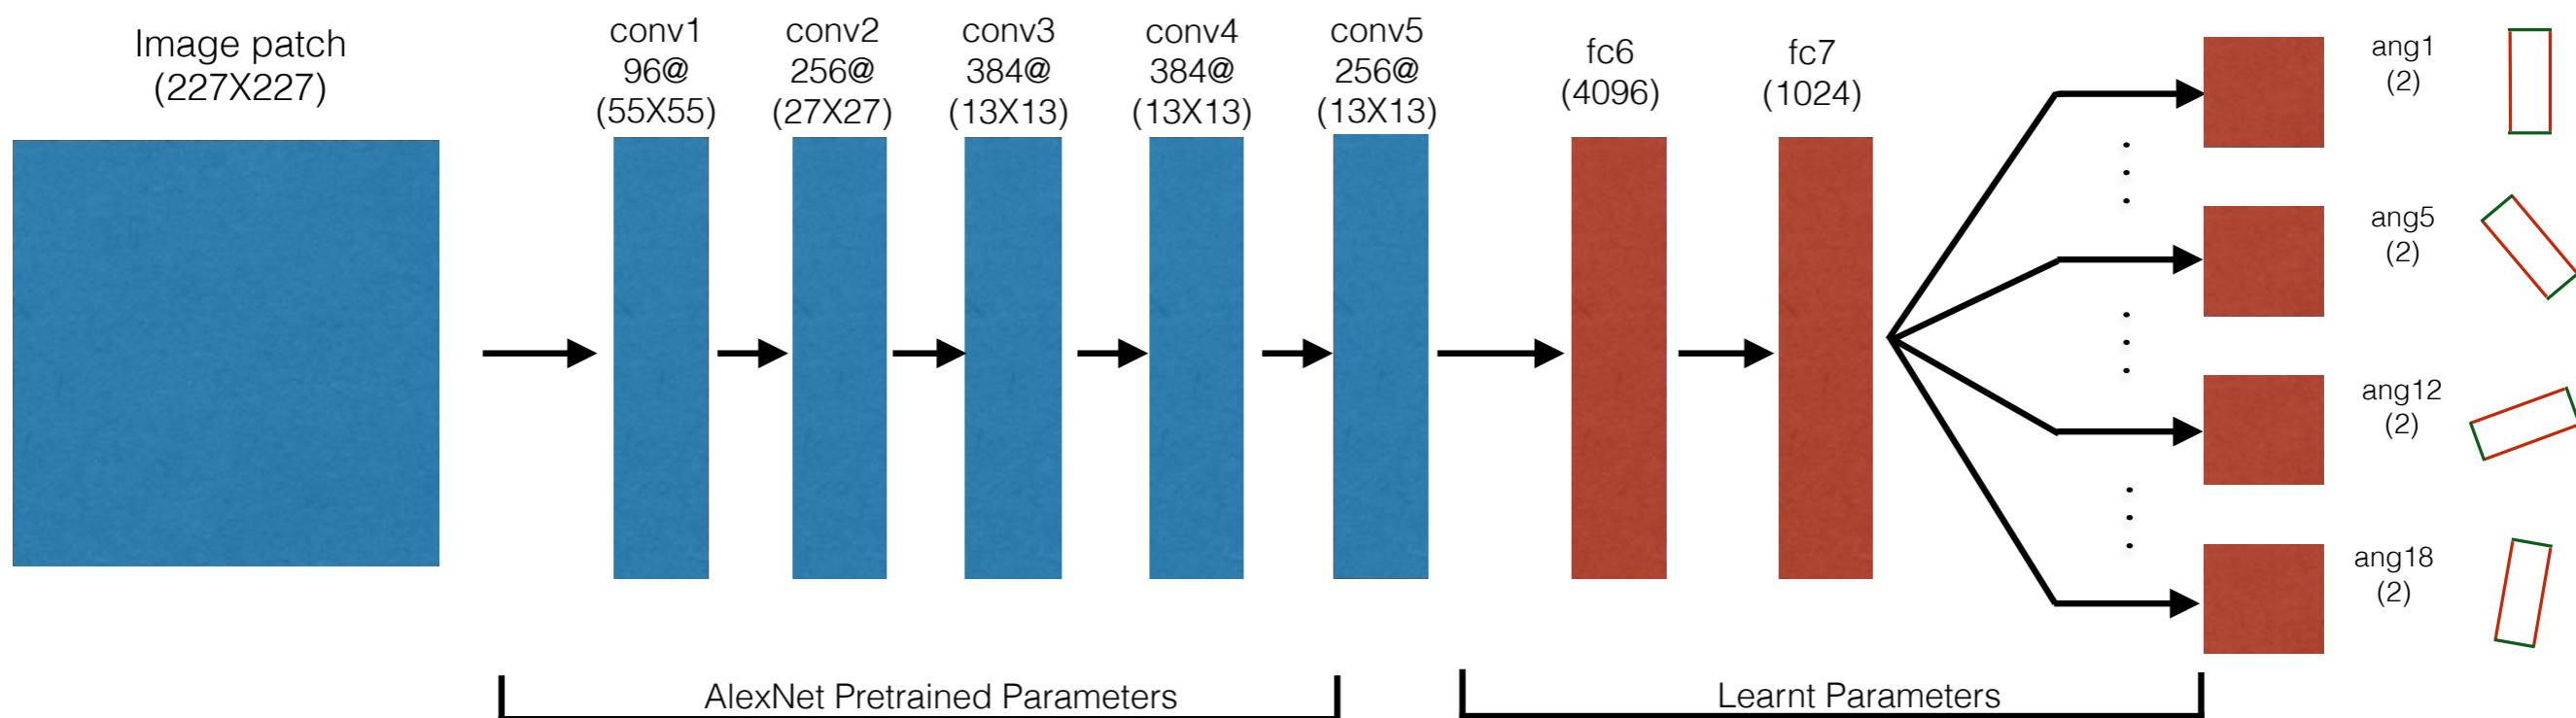
Positive Grasp Patches



Negative Grasp Patches



# Self-supervised Grasp Affordances

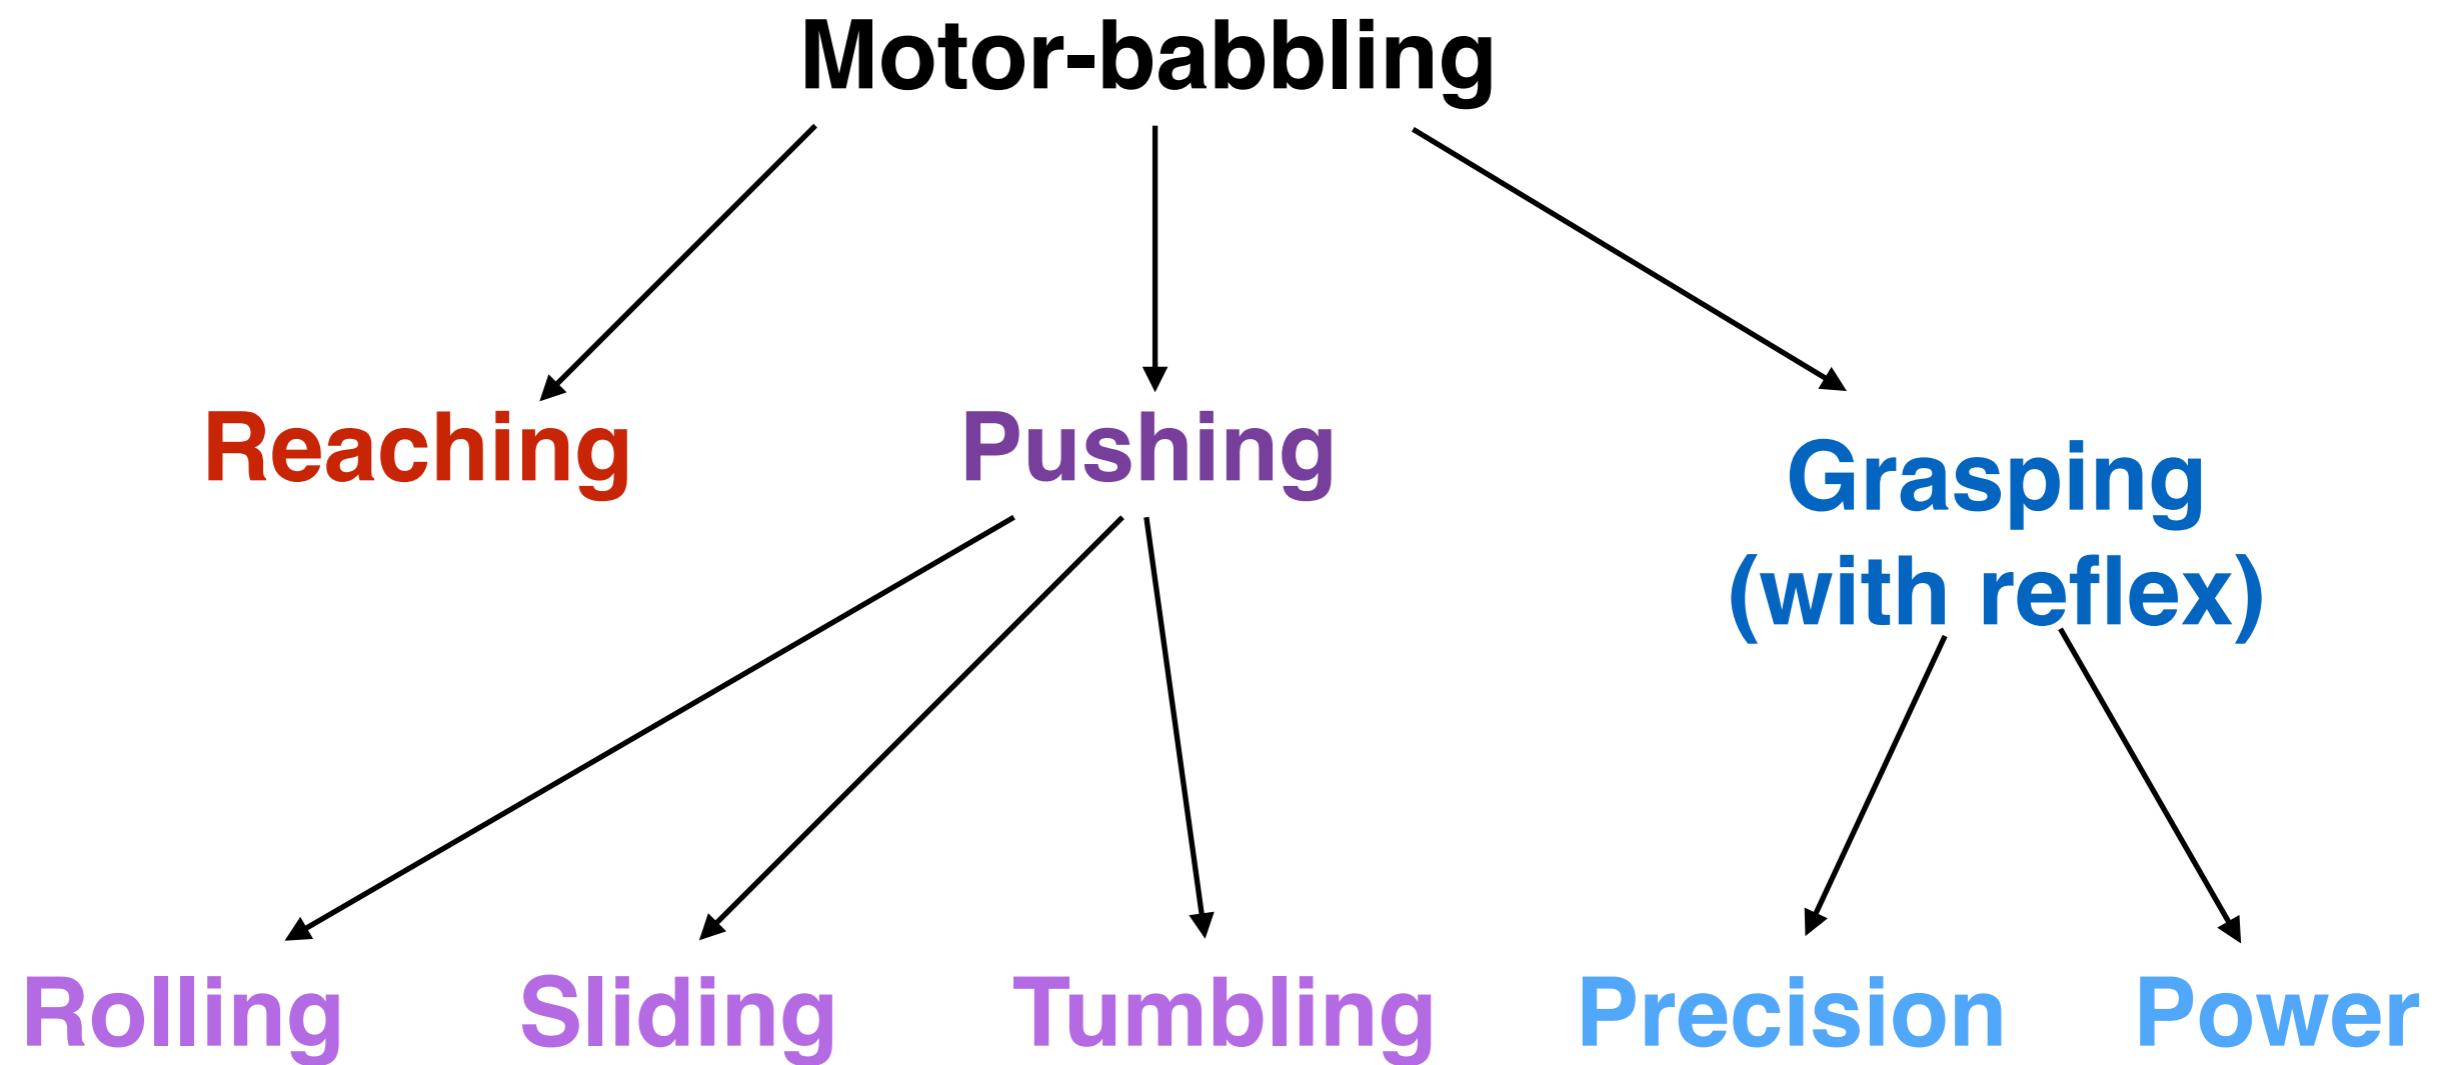


# Self-supervised Grasp Affordances



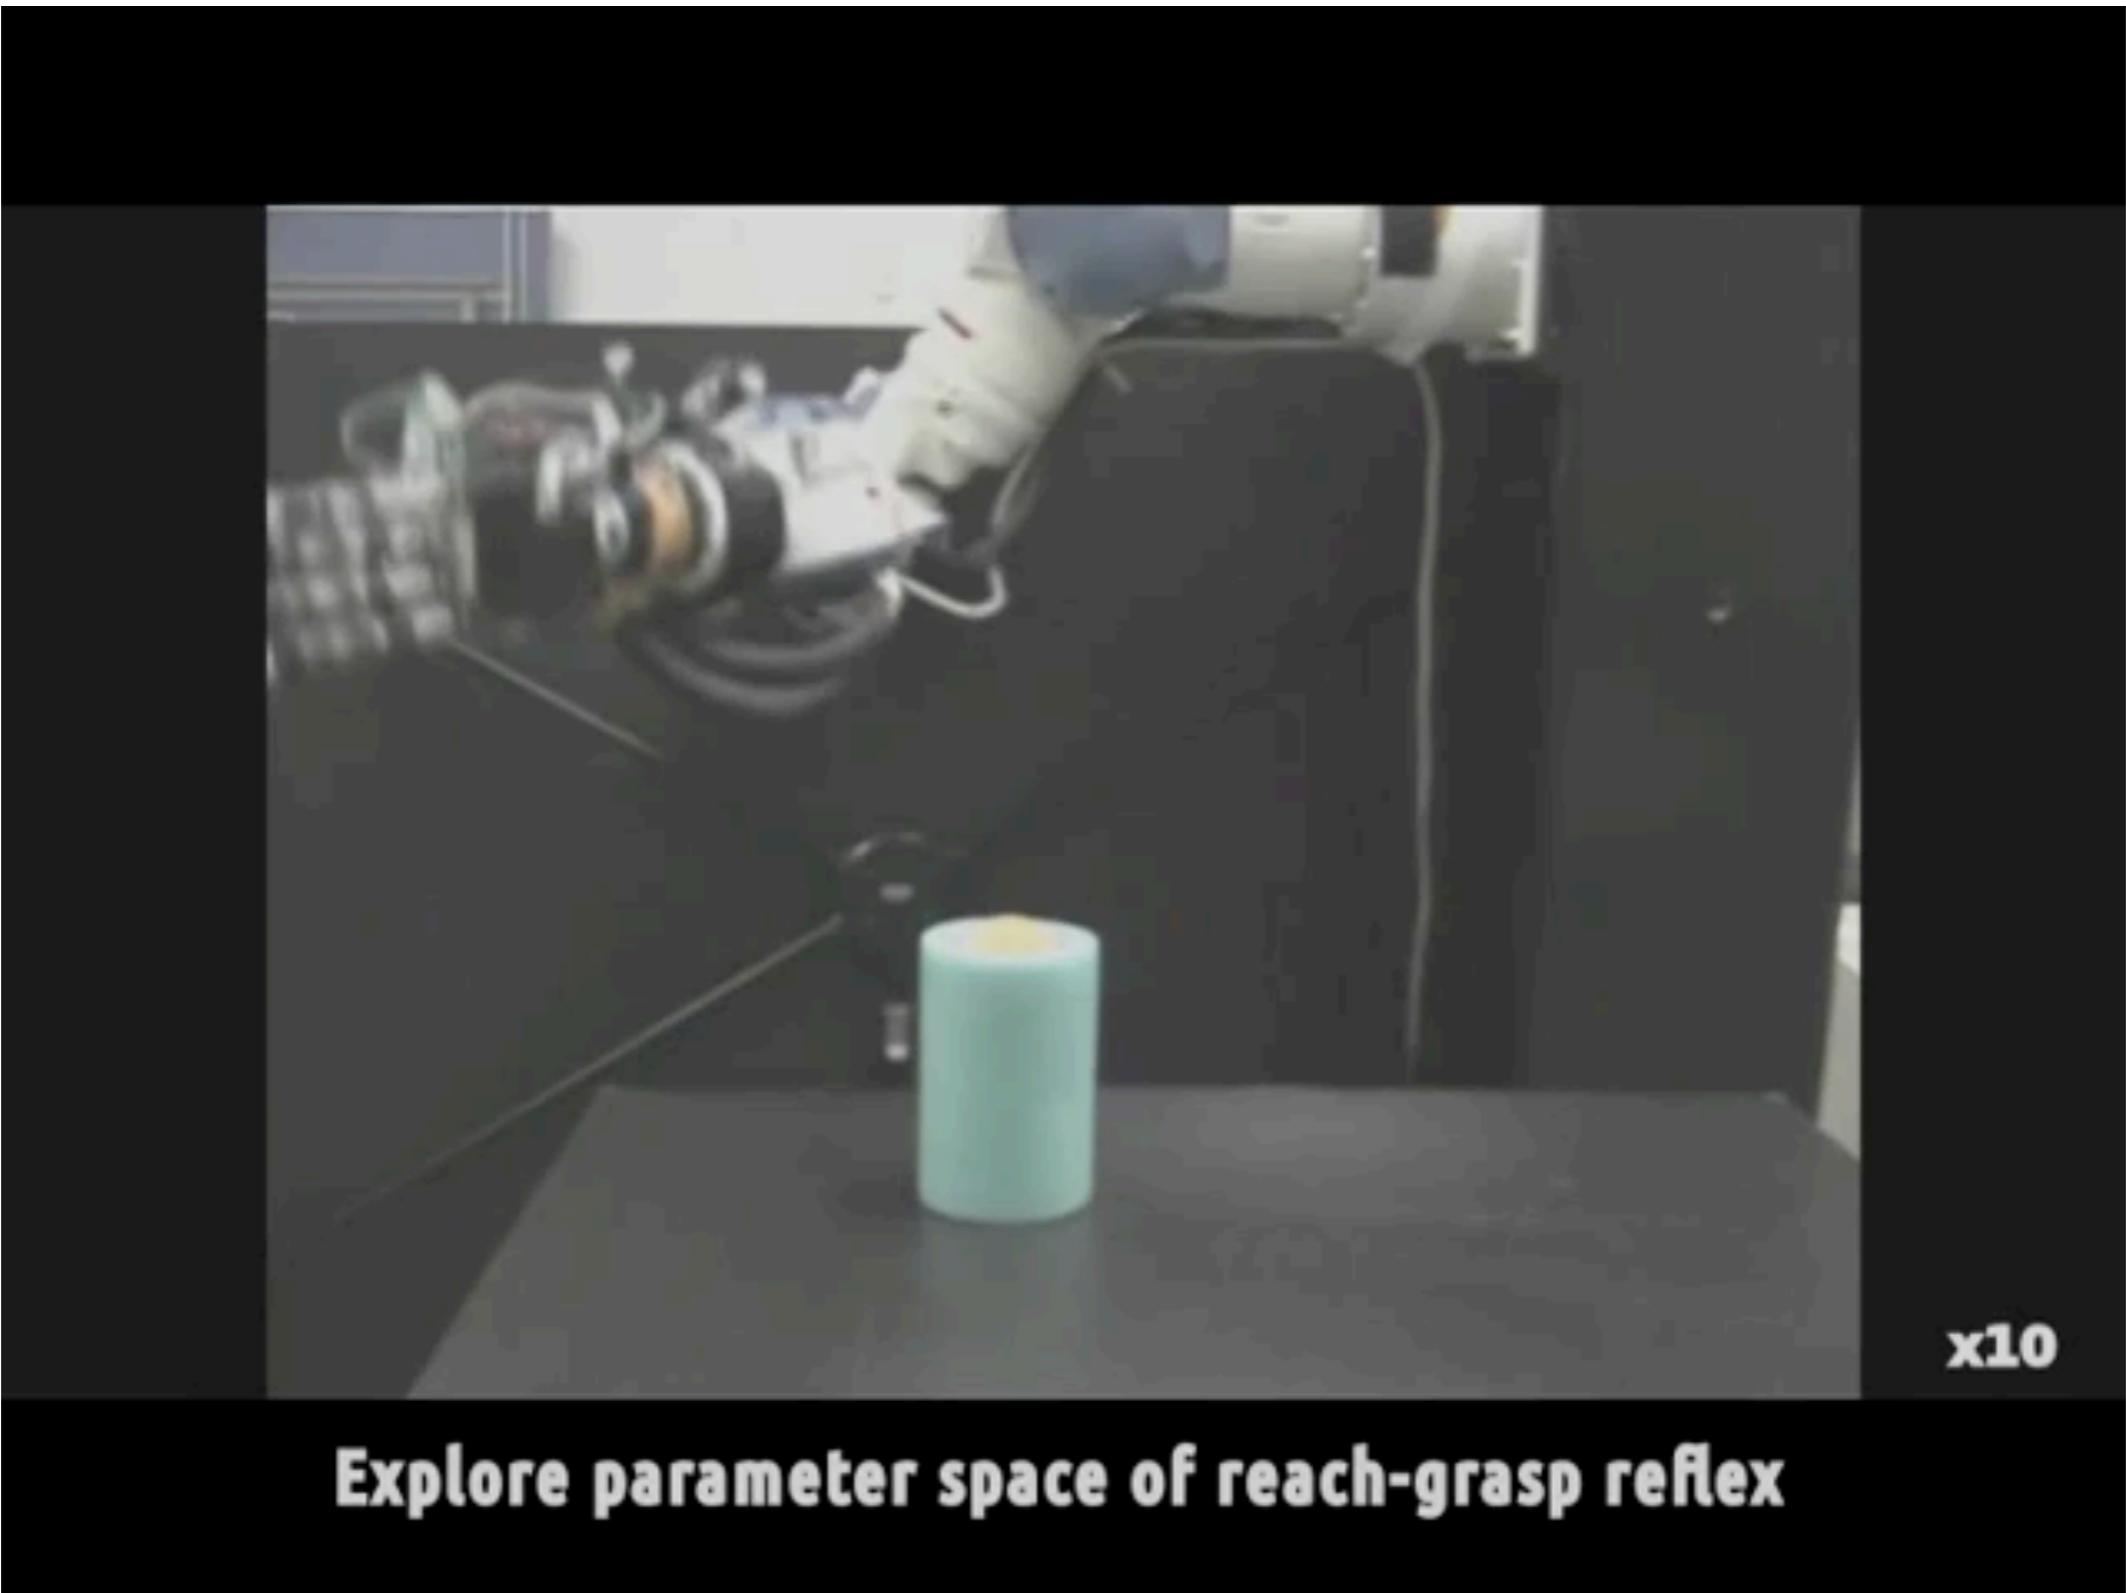
# Developmental Learning Example

- Interactions often lead to distinct salient effects
- Can cluster effects to create specific distinct skills



- Learn states that afford each skill (aka preconditions)

# Developmental Learning Example



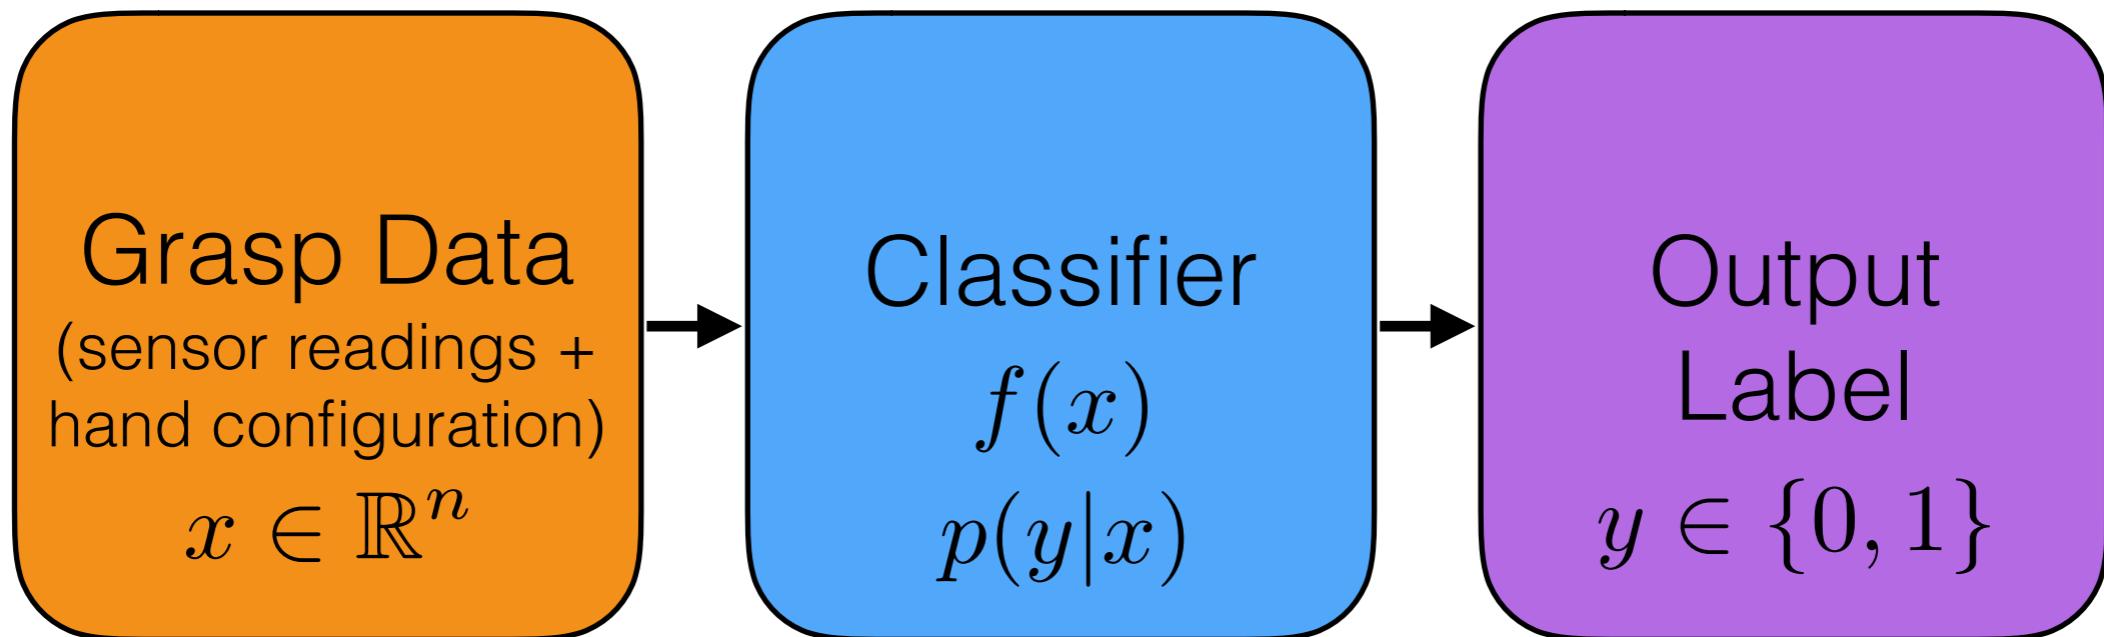
---

# **Logistic Regression**

---

# Classifying Grasps

- **Classification** problem from machine learning viewpoint



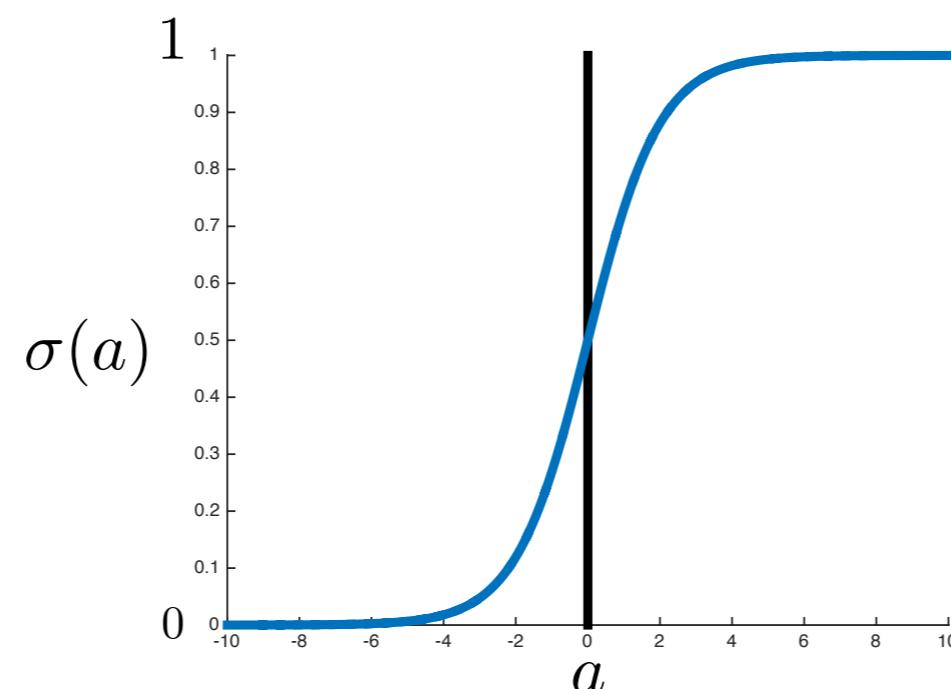
- **Discriminant function** defines decision boundary  
$$f(x) = 0$$
- **Probabilistic classifier** models posterior distribution  
$$p(y|x)$$

# Logistic Regression

- Model posterior distribution using **sigmoid** function

$$p(y = 1|x) = \sigma(\theta^T x) = \frac{1}{1 + \exp(-\theta^T x)}$$

$$p(y = 0|x) = 1 - \sigma(\theta^T x)$$



- Assume  $x$  contains a constant bias term of 1
- Need to learn the parameters  $\theta$  from training data

# Logistic Regression Training

- Training data consists of  $N$  labeled data points

$$y^{(1)}, y^{(2)}, \dots, y^{(N)} \quad x^{(1)}, x^{(2)}, \dots, x^{(N)}$$

- Likelihood of the model parameters

$$L(\theta) = \prod_{i=1}^N \sigma(\theta^T x^{(i)})^{y^{(i)}} (1 - \sigma(\theta^T x^{(i)}))^{1-y^{(i)}}$$

- Negative log likelihood error function (cross entropy)

$$E(\theta) = -\log(L(\theta)) = -\sum_{i=1}^N y^{(i)} \log(\sigma(\theta^T x^{(i)})) + (1 - y^{(i)}) \log(1 - \sigma(\theta^T x^{(i)}))$$

- Optimize parameters using gradient decent

$$\nabla E(\theta) = \sum_{i=1}^N (\sigma(\theta^T x^{(i)}) - y^{(i)}) x^{(i)} \quad \theta_{\text{new}} = \theta_{\text{old}} - \beta \nabla E(\theta_{\text{old}})$$

learning rate



# Logistic Regression Example

---

- Compute the updated parameters for the following data

$$x^{(1)} \begin{bmatrix} 0 \\ 2 \\ 1 \end{bmatrix}, y^{(1)} = 1$$

$$x^{(2)} \begin{bmatrix} 1 \\ 0.5 \\ 1 \end{bmatrix}, y^{(2)} = 0$$

$$x^{(3)} \begin{bmatrix} -1 \\ 0 \\ 1 \end{bmatrix}, y^{(3)} = 1$$

$$\theta_{\text{old}} = [ \ 1 \ 1 \ -1 \ ]^T, \beta = 0.1$$

$$x^{(1)} = \begin{bmatrix} 0 \\ 2 \\ 1 \end{bmatrix}, y^{(1)} = 1 \quad x^{(2)} = \begin{bmatrix} 1 \\ 0.5 \\ 1 \end{bmatrix}, y^{(2)} = 0 \quad x^{(3)} = \begin{bmatrix} -1 \\ 0 \\ 1 \end{bmatrix}, y^{(3)} = 1$$


---

$$\theta_{\text{old}} = [1 \ 1 \ -1]^T, \beta = 0.1$$

$$\nabla E(\theta) = \sum_{i=1}^N (\sigma(\theta^T x^{(i)}) - y^{(i)}) x^{(i)}$$

- **Compute current predictions**

$$\sigma(0 + 2 - 1) = 0.7311$$

$$\sigma(1 + 0.5 - 1) = 0.6225$$

$$\sigma(-1 + 0 - 1) = 0.1192$$

- **Compute errors**

$$0.7311 - 1 = -0.2689$$

$$0.6225 - 0 = 0.6225$$

$$0.1192 - 1 = -0.8808$$

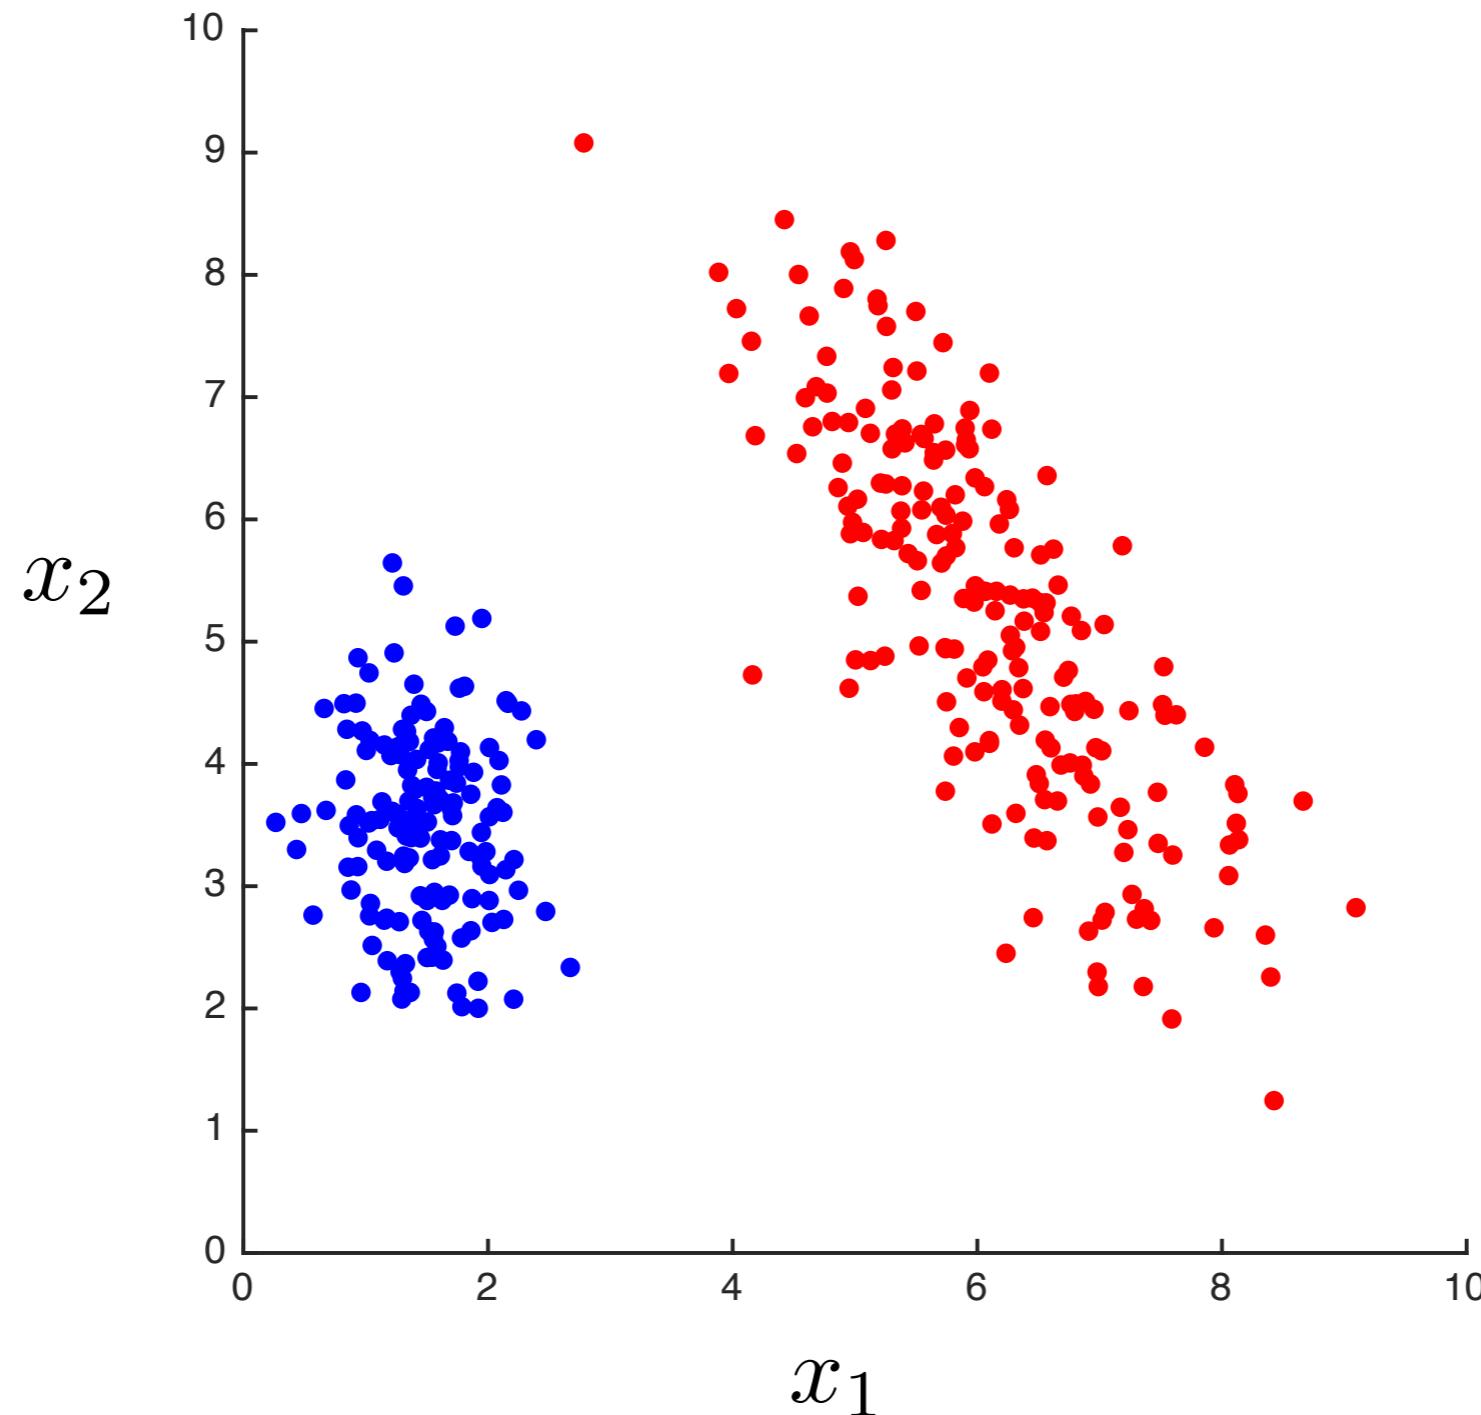
- **Compute gradient and update parameters**

$$\nabla E = -0.2689 \begin{bmatrix} 0 \\ 2 \\ 1 \end{bmatrix} + 0.6225 \begin{bmatrix} 1 \\ 0.5 \\ 1 \end{bmatrix} - 0.8808 \begin{bmatrix} -1 \\ 0 \\ 1 \end{bmatrix} = \begin{bmatrix} 1.5033 \\ -0.2267 \\ -0.5273 \end{bmatrix}$$

$$\theta_{\text{new}} = \begin{bmatrix} 1 \\ 1 \\ -1 \end{bmatrix} - 0.1 \begin{bmatrix} 1.5033 \\ -0.2267 \\ -0.5273 \end{bmatrix} = \begin{bmatrix} 0.8497 \\ 1.0227 \\ -0.9473 \end{bmatrix}$$

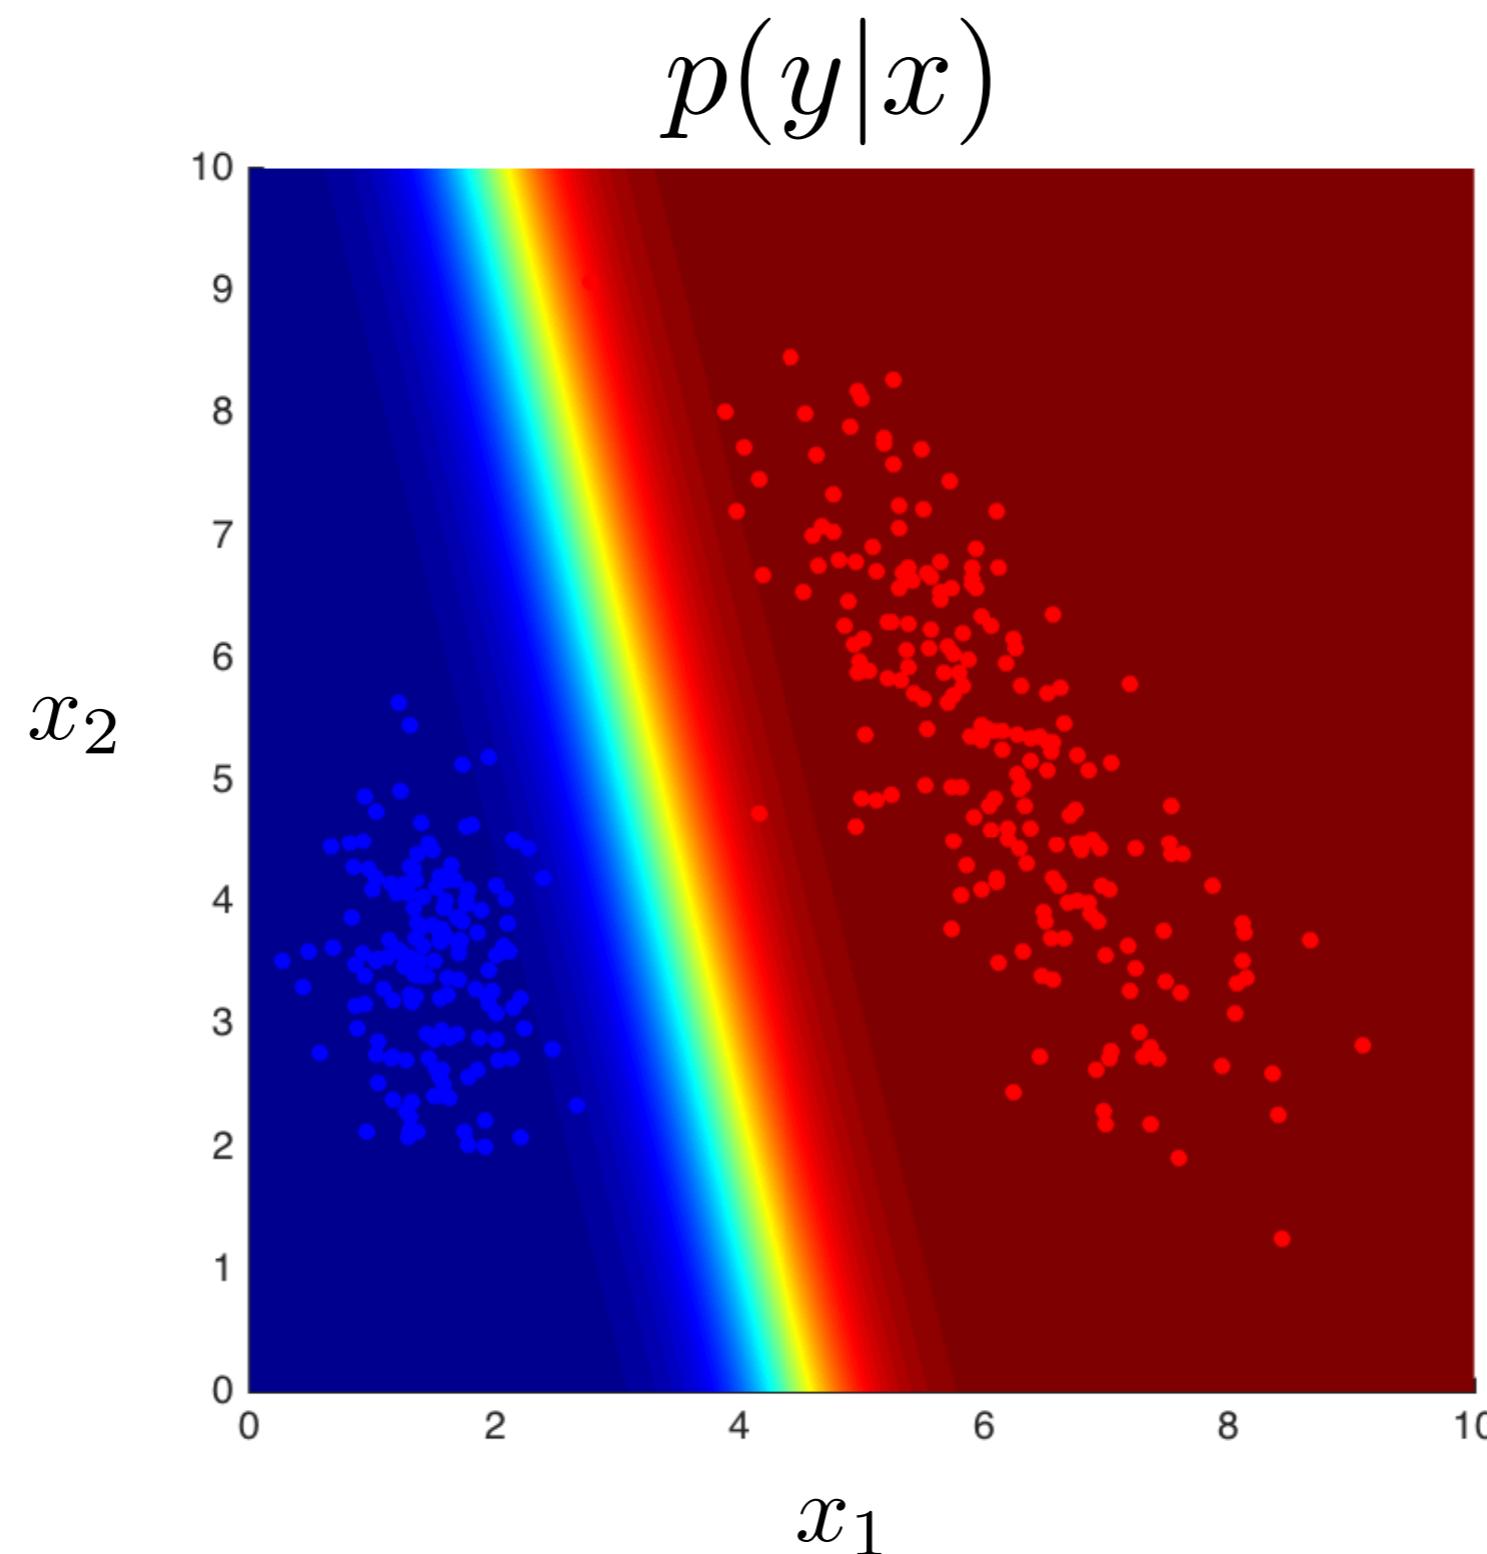
# Logistic Regression Example

---



# Logistic Regression Example

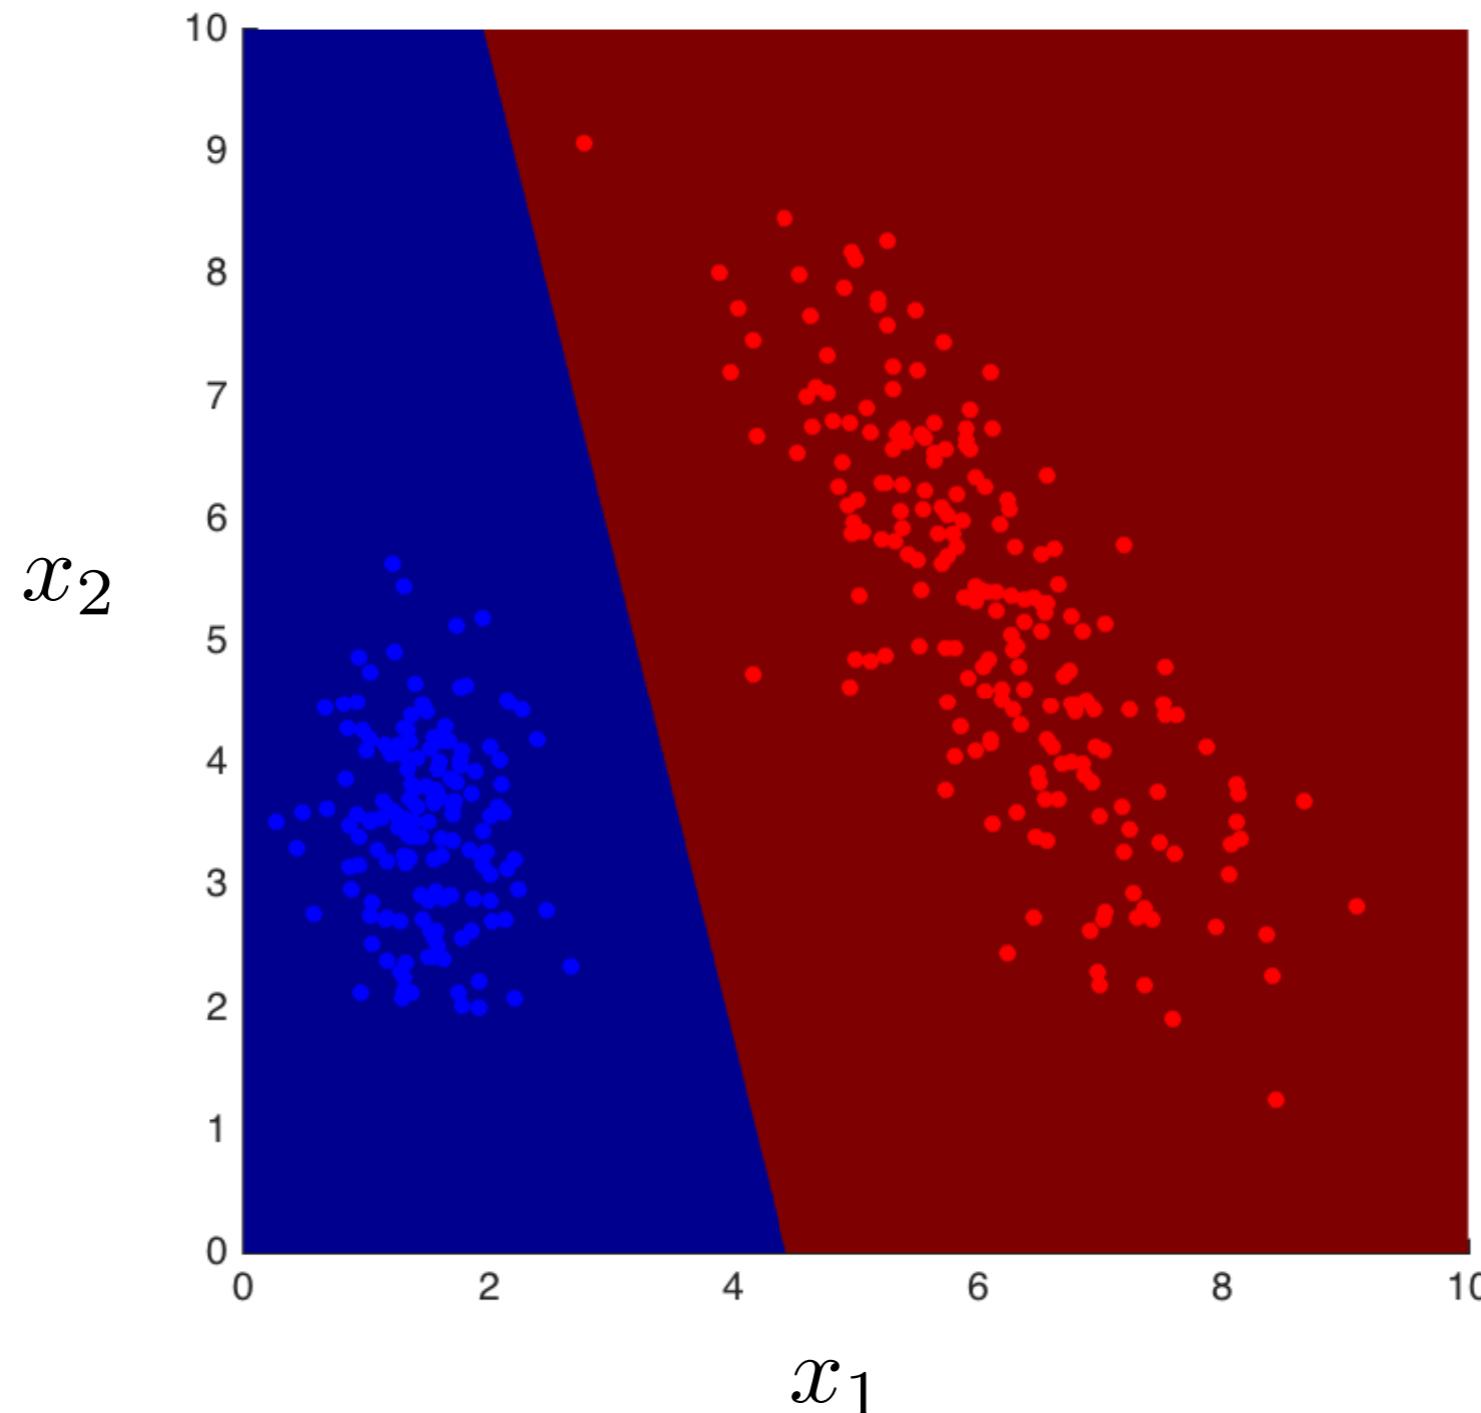
---



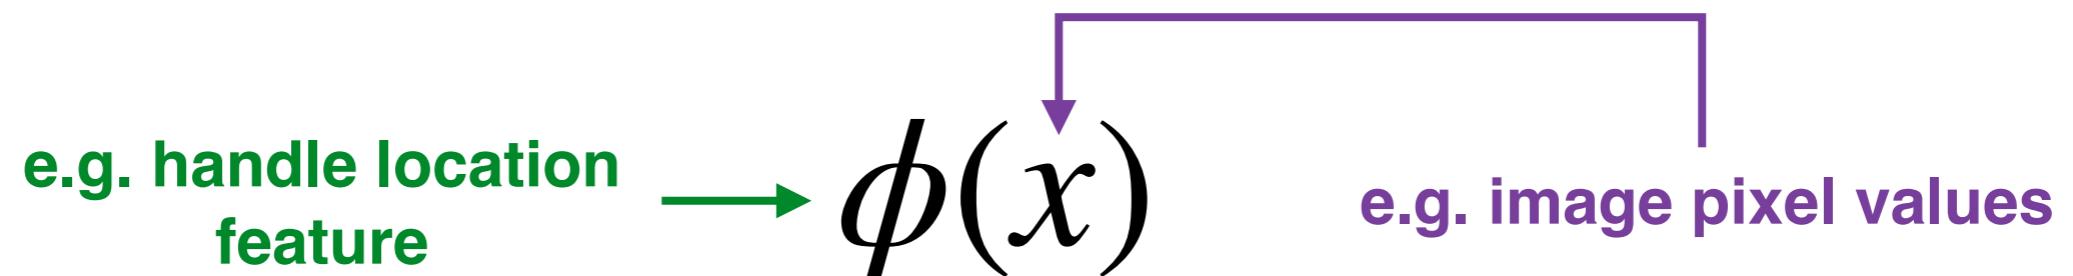
# Logistic Regression Example

---

## Decision Boundary



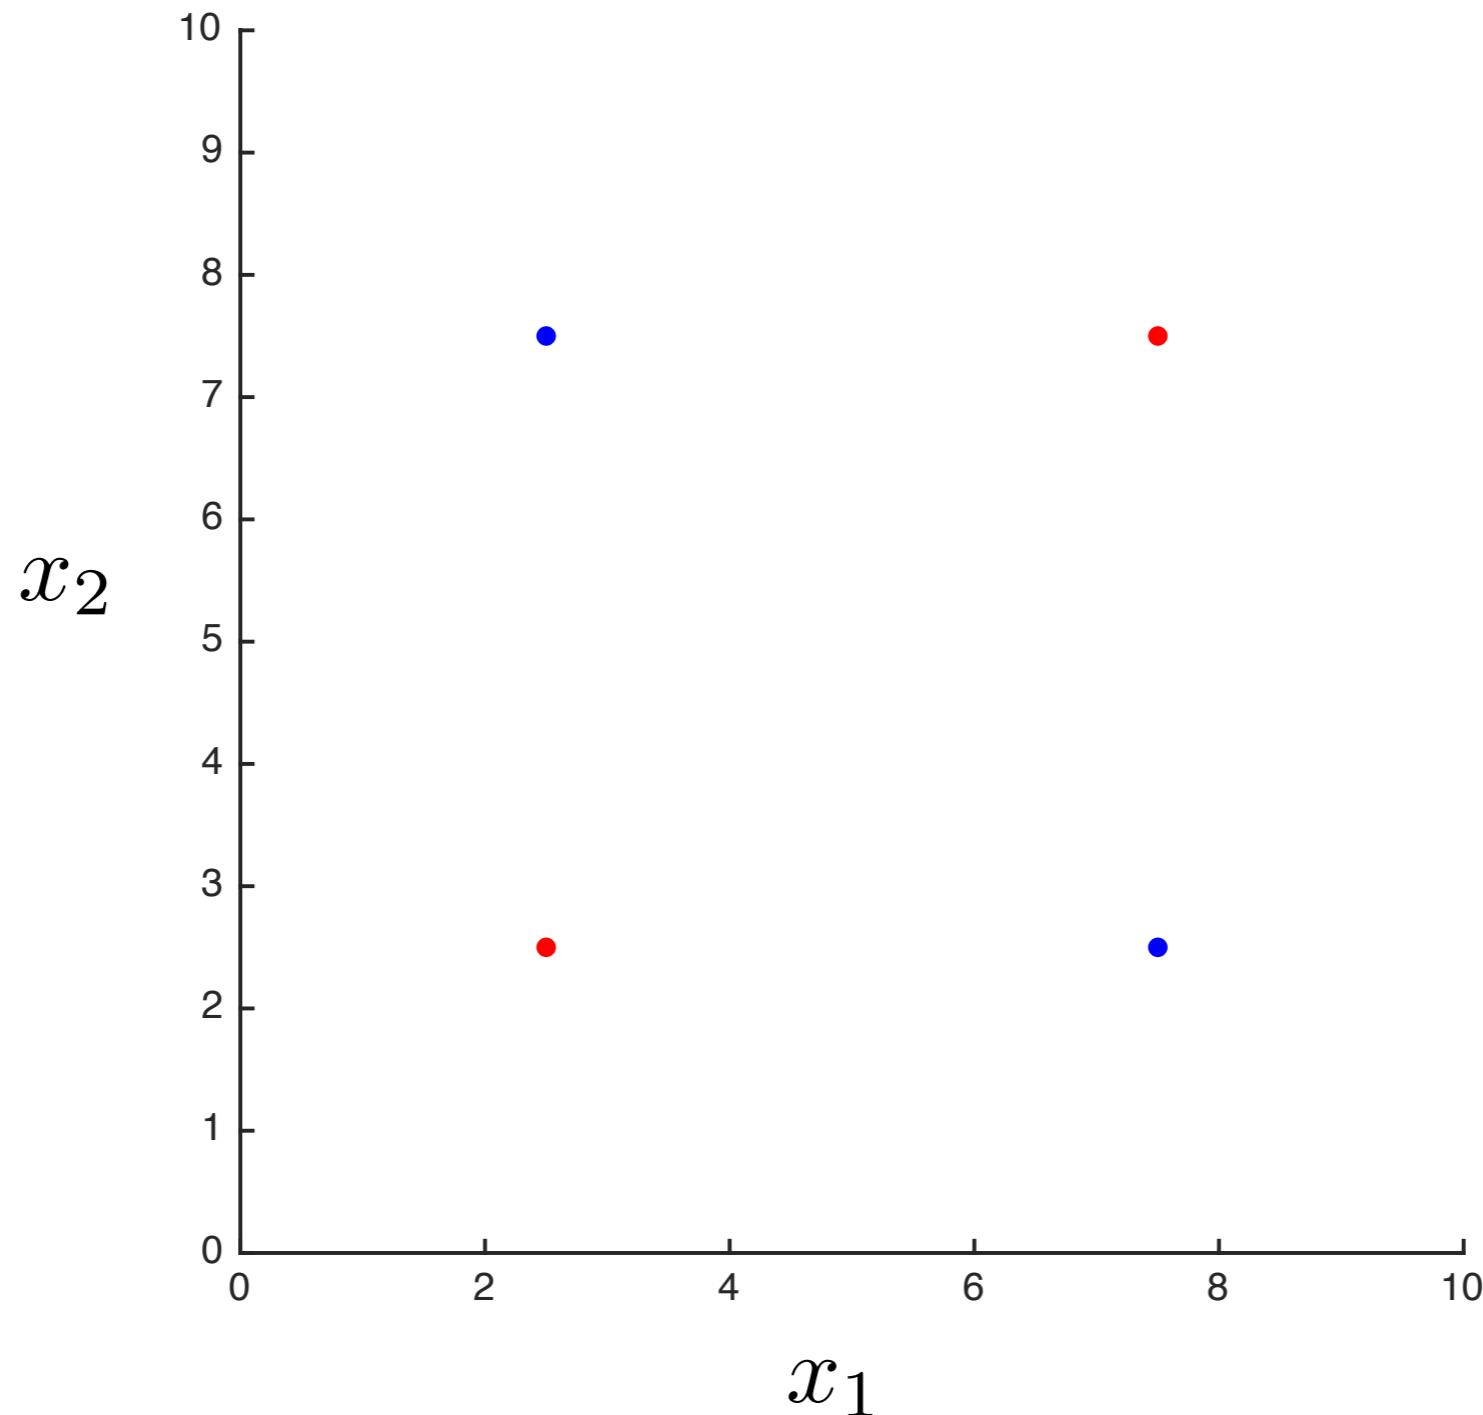
- A key challenge is finding a suitable **input representation**
- In practice we often talk about **features** of the input



- Logistic regression same as before with  $\phi(x)$  instead of  $x$ 
  - ▶ Nonlinear features allow for nonlinear boundaries in  $x$  space
  - ▶ Neural networks can be seen as learning features

# Neural Network XOR Example

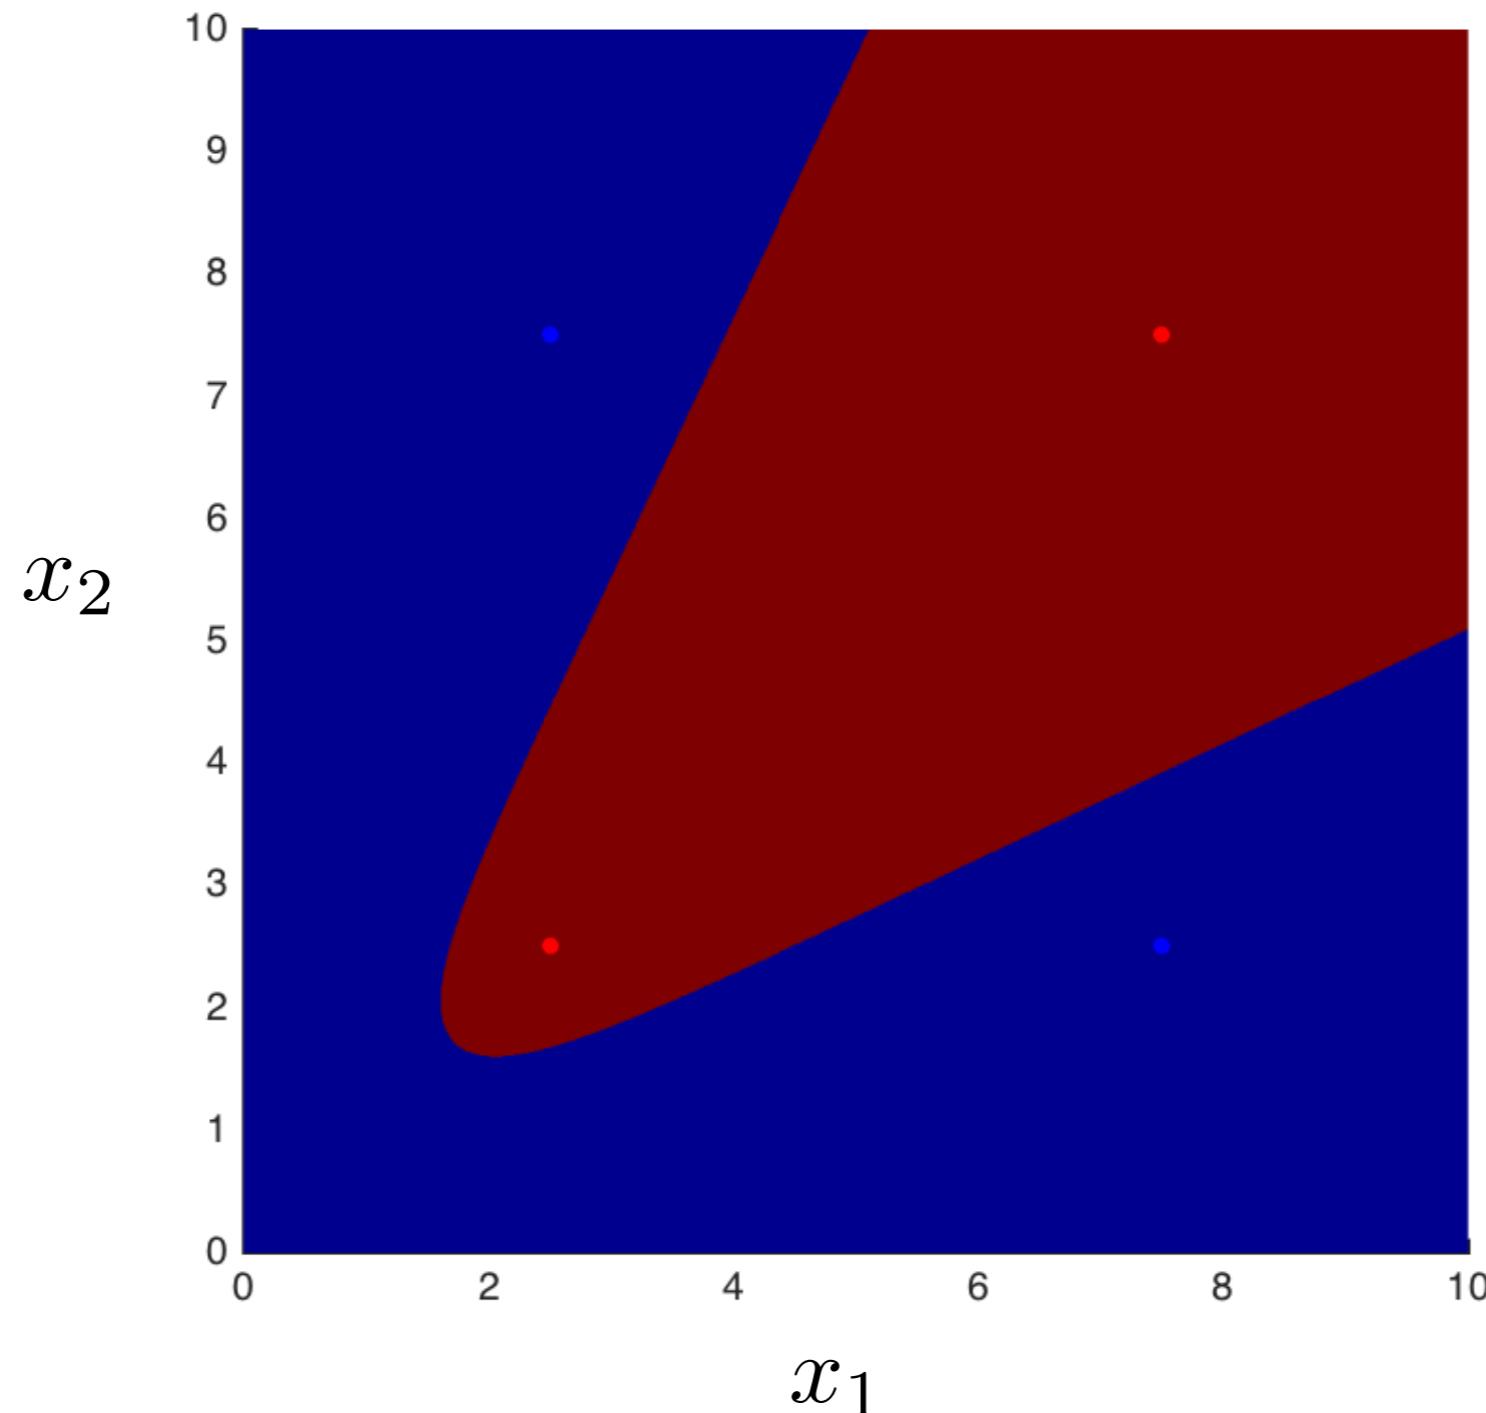
---



# Neural Network XOR Example

---

## Decision Boundary



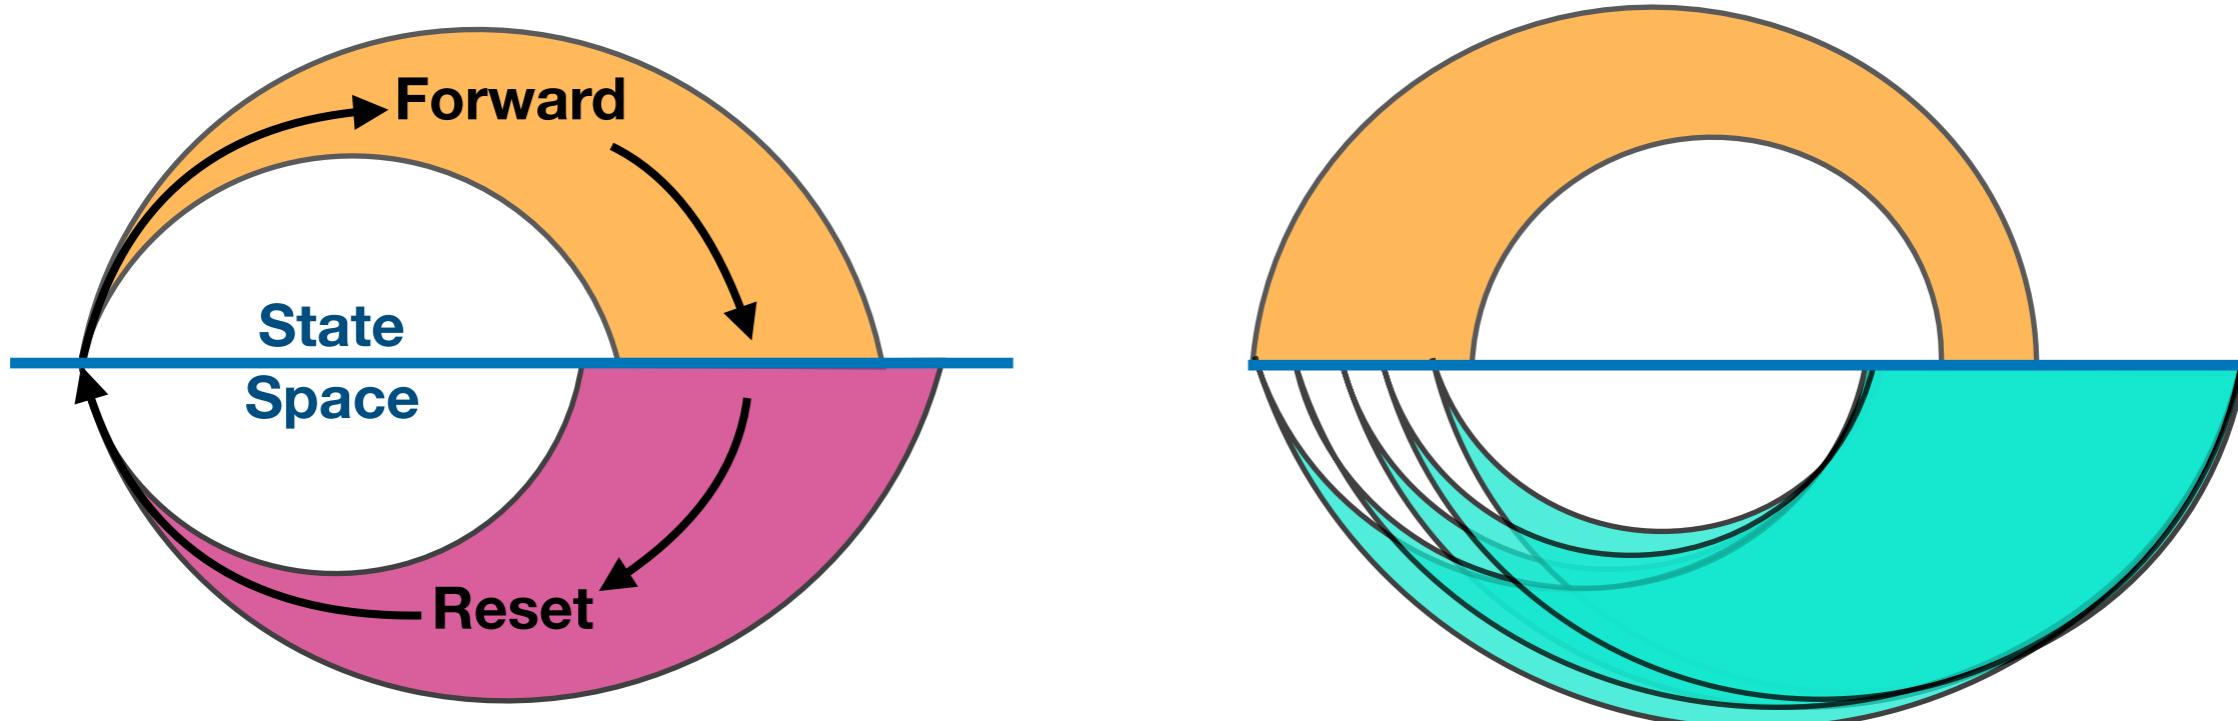
---

# **Autonomous Resetting**

---

# Resetting Environment

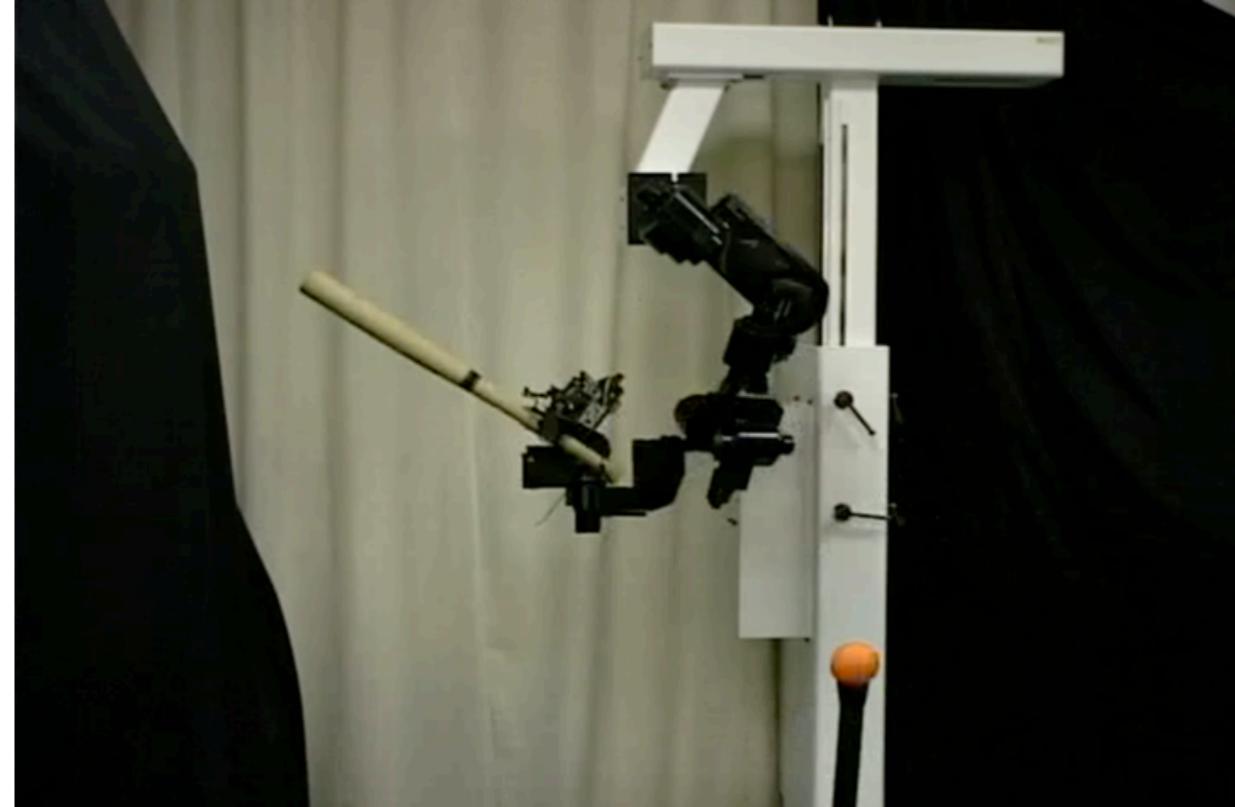
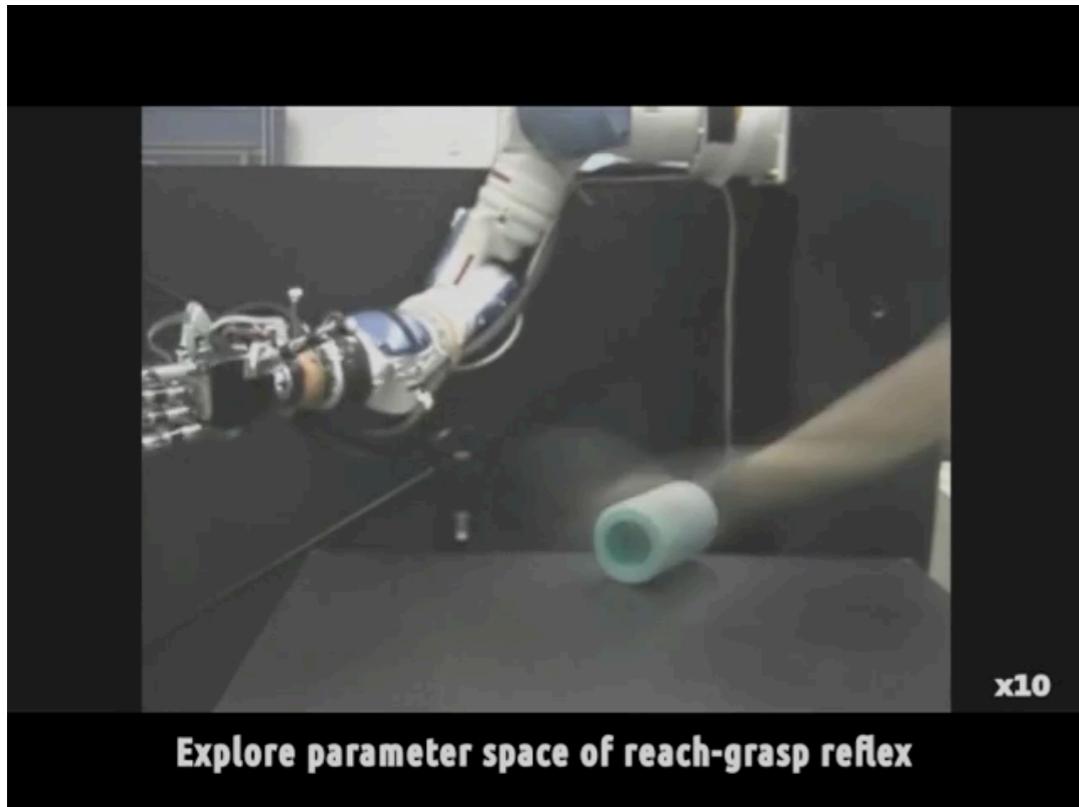
- Learning skills and models requires training data
- Often need to **reset** objects between trials



- Simulations allow for simple resetting
- Want to collect **physical** data from real world

# Human in the Loop

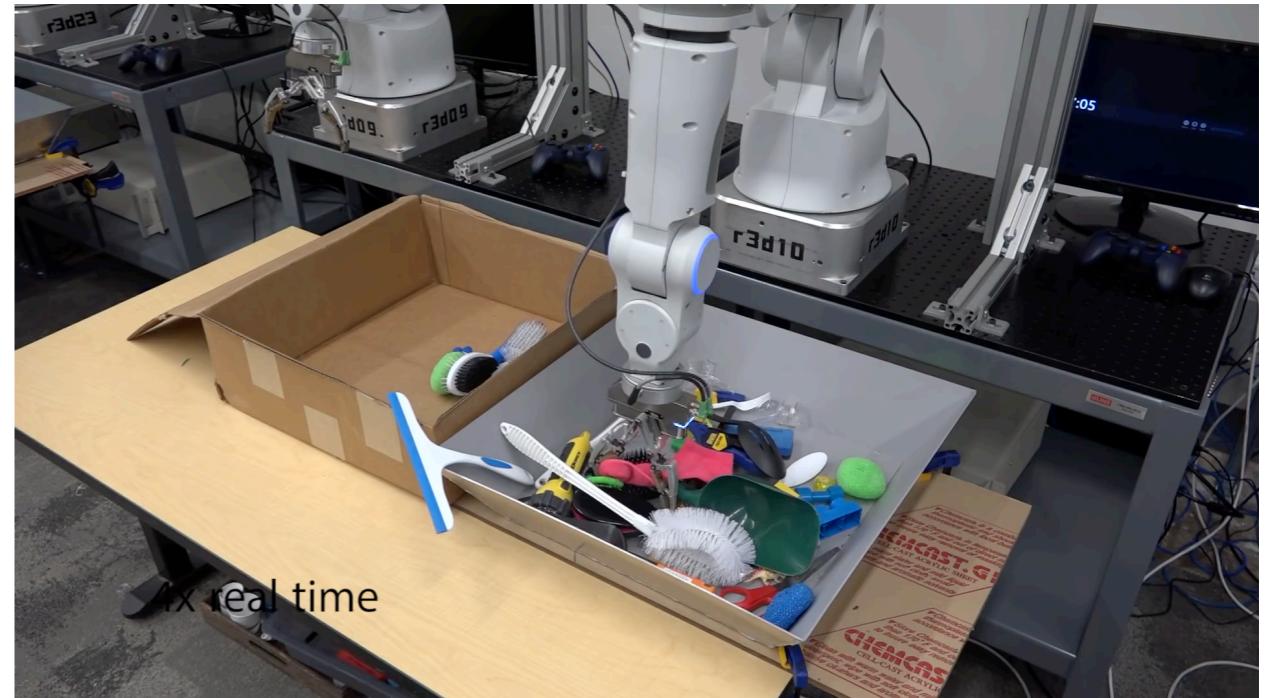
- Have a human reset the environment between trials



- ▶ Tedious and limits robot's overall autonomy
- ▶ Easy to explore a wide range of structured situations

# Working in Clutter

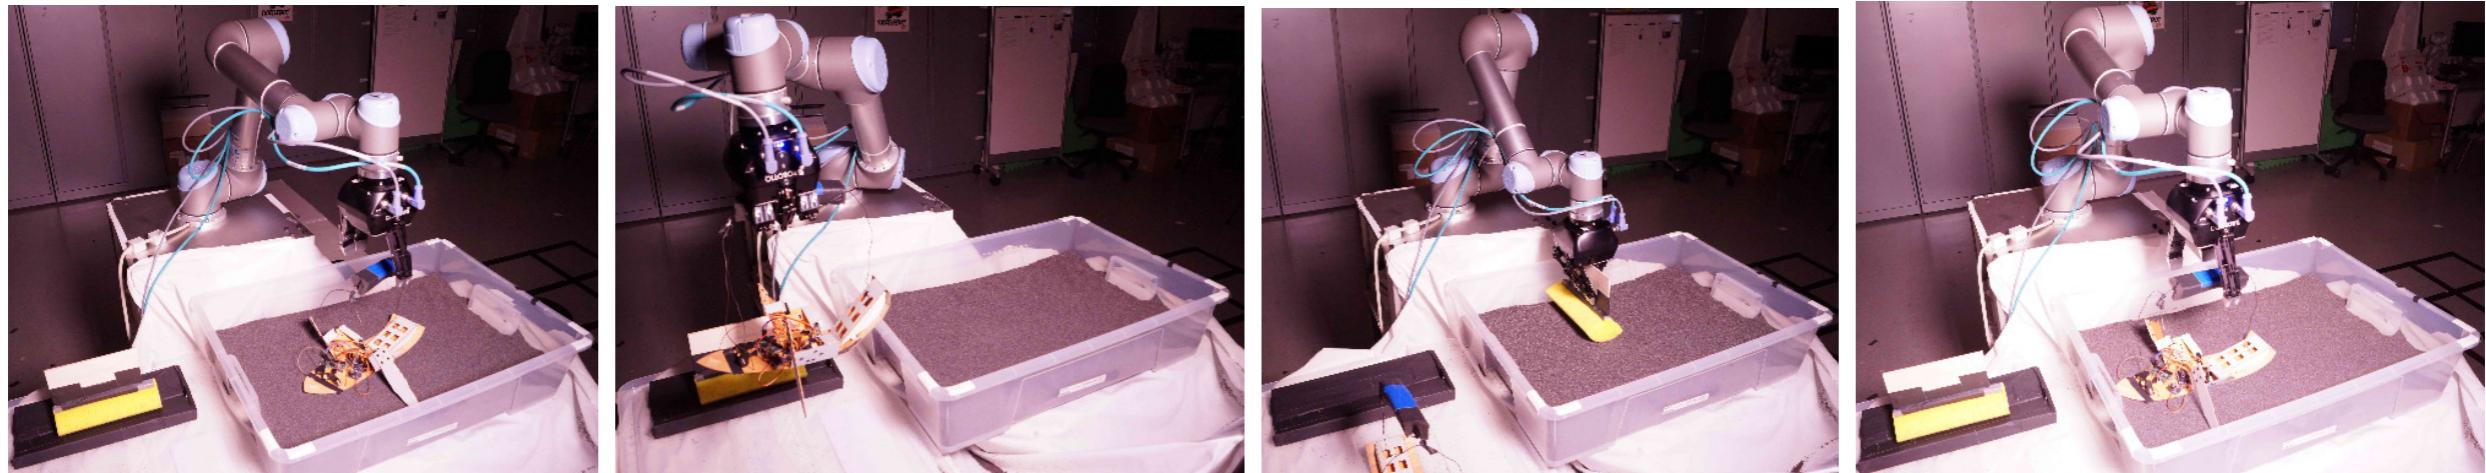
- Work in clutter to easily create diverse environments



- ▶ Create new cluttered state by dropping objects
- ▶ Target rich environments for simple actions
- ▶ May not reflect real world scenes

# Resetting Robot Actions

- Use robot to reset environment between trials



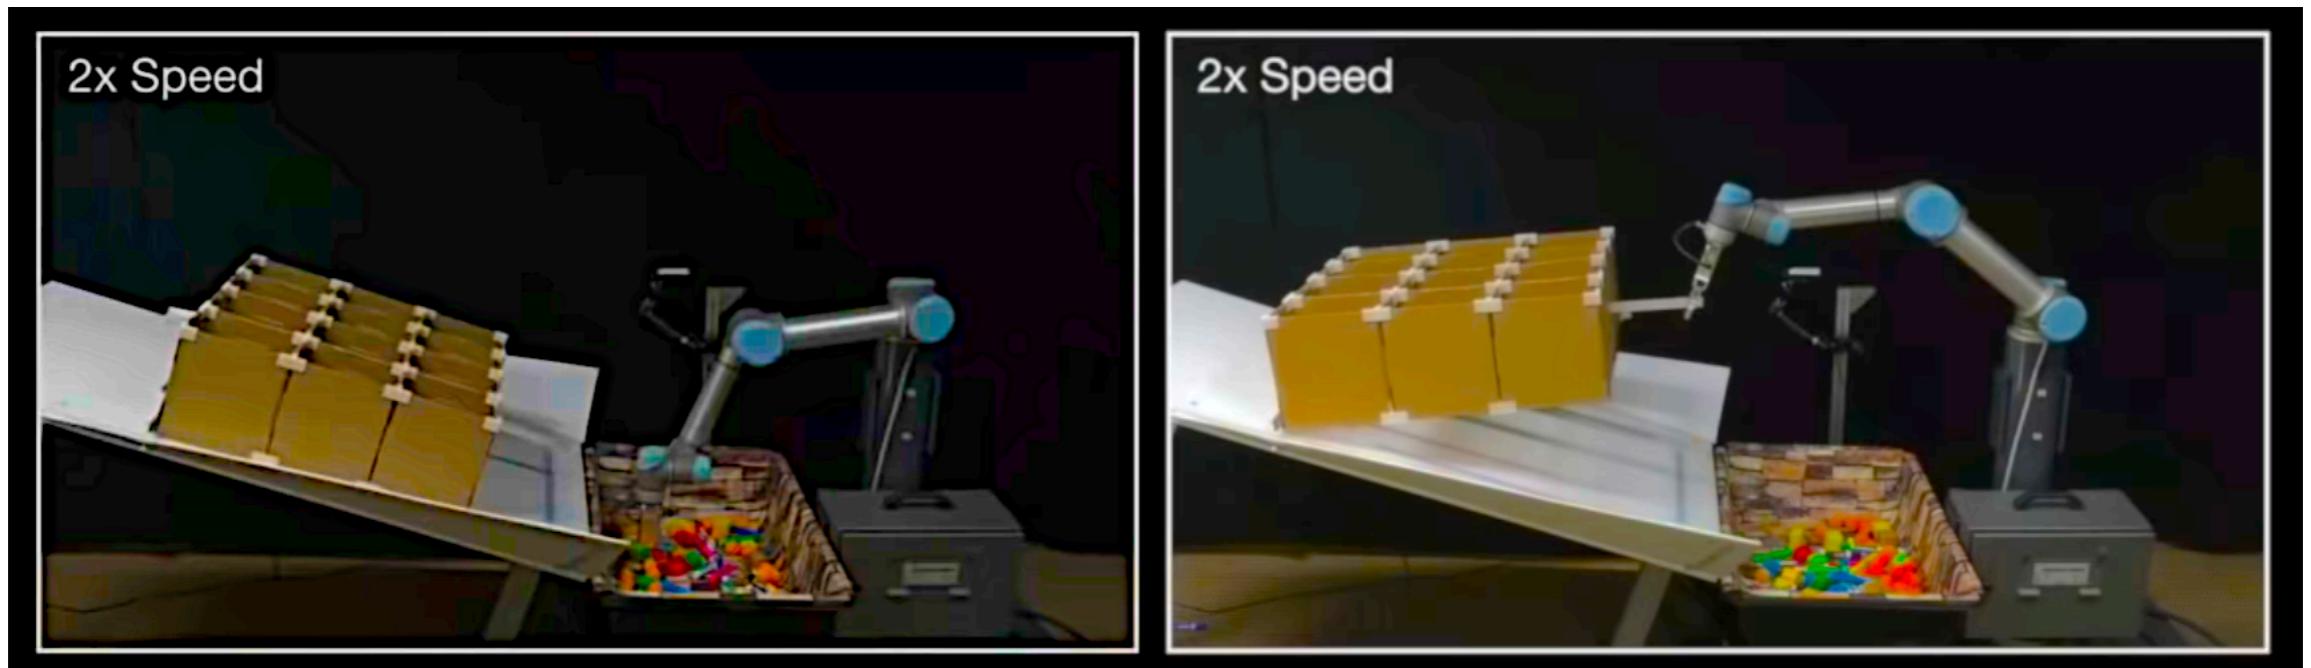
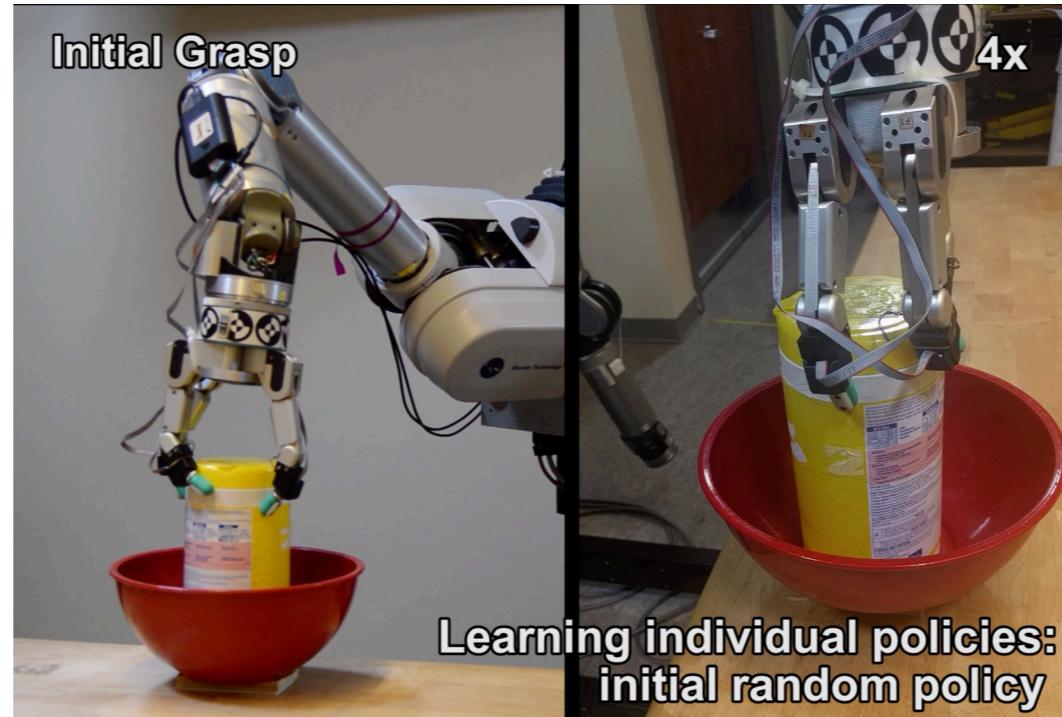
- ▶ Use pre-programmed actions to reset environment



- ▶ Can augment objects to allow for more interactions, e.g., pulling

# Resetting Mechanisms

- Create structures in environment to help reset trials



- Robots can learn models from **experience**
  - ▶ Perform actions and observe the resulting effects
- **Interactive perception**
  - ▶ Observe latent information by interacting with objects
- **Transition models**
  - ▶ Predict the effects of different actions
- **Affordances**
  - ▶ Determine if an object/state affords a desired skill
- **Autonomous resetting** allows for easier data collection

---

Questions?